

COEN 175

Lecture 15: Stack-Allocated Variables

Parameter Passing

- Who evaluates the arguments to the called function?
 - The caller evaluates them since they are passed by value.
- How are arguments communicated to the callee?
 - Some arguments may be placed in registers.
 - Other arguments will be placed on the stack.
- Who evaluates the return value?
 - The callee evaluates the return value.
- How is the return value communicated to the caller?
 - The value is typically placed in a register (assuming it fits).

Calling Conventions

- The **calling conventions** of the system dictate how the registers and the stack are used.
 - Which registers are used to pass arguments?
 - What if an argument is too large for a single register?
 - Where are extra arguments placed on the stack?
 - Which register is used to return a value?
 - Where is a return value placed if it does not fit in a register?
 - Which registers can be modified by the callee?
- These conventions depend on the OS used.
 - Unix and Microsoft systems have different conventions.


Intel Register Names

- The original 16-bit Intel register names:
 - AX – accumulator
 - BX – base register
 - CX – count register
 - DX – data register
 - SI – source index
 - DI – destination index
- The 32-bit versions are called EAX, EBX, etc.
- The 64-bit versions are called RAX, RBX, etc.
 - The additional 8 registers are simply R8–R15.
 - The 32-bit versions of these are called R8d–R15d.


Linux Calling Conventions

- No arguments are passed in registers.
- All arguments are placed on the stack.
 - Arguments are pushed on the stack from **right to left**.
 - Arguments must be a multiple of 4 bytes.
 - The stack must be kept aligned on a 4-byte boundary.
(Note that OS X requires 16-byte alignment!)
 - The first argument is +8 bytes from the base pointer.
- An integer return value is placed in EAX.
 - A floating value is returned differently using the floating-point co-processor stack.


Stack-Based Parameters



Before call to
`foo(arg1, arg2)`




Push rightmost arg



Push next argument

Stack-Based Parameters



Referencing Locals

- Local variables are stored within the stack frame, which could be located anywhere in memory.
- Thus, we cannot use absolute or fixed addressing.
- Instead, we must use relative addressing.
 - Also called base-offset addressing.
- Relative to which base register? FP or SP?
 - SP might move during function execution.
 - FP will not move during function execution.
 - So, FP is a better choice.

Assigning Offsets

- Local variables will be referenced via negative offsets from the frame pointer.
- For simplicity, local variables are typically laid out within the frame in declaration order:
 - Keep track of the last offset in a global variable offset.
 - Assign the next local variable, x, its offset as follows:

```
offset = offset - x.size
x.offset = offset
```
- For example, if the first variable requires 4 bytes, it would be at offset -4. If it required 8 bytes, it would be at offset -8.

Example: Offsets

- Assign offsets to each variable:

```
int q, a[10];           // none: q and a are static


int f(int i, int j) {    // i = +8, j = +12
    int x, *p;           // x = -4, p = -8
    double z;             // z = -16
    char a[4], *s;        // a = -20, s = -24
}
```

Example: Offsets

- Assign offsets to each variable:

```
int q, a[10];
```

```
int f(int i, int j) {
    int x, *p;
    double z;
    char a[4], *s;
}
```



- What if `f()` were to call another function?
- What would happen to any registers in use?

Register Usage

- There is only one set of registers that must be shared amongst all functions.
- Most calling conventions divide up the registers:
 - **Caller-saved** – the register's value must be saved on the stack before a function call and restored after the call;
 - **Callee-saved** – the register's value must be saved before it is first used and restored before returning;
 - Any argument and return registers must be caller-saved.
- Intel 32-bit convention: EAX, ECX, and EDX are caller-saved; EBX, ESI, and EDI are callee-saved.

Another Example: Offsets

- Assign offsets to each variable:

```
int f(int i, int j) {      // i = +8, j = +12
    char x, y, *p;        // x = -1, y = -2, p = -6
    double z;              // z = -14
    char a[5], *s;         // a = -19, s = -23
}
```

- Some of our objects are not **aligned**.
 - Not all architectures require alignment.
 - An unaligned memory access will trigger a **bus error**.
 - Intel does not require alignment, but aligned accesses run faster as unaligned accesses require multiple fetches.

Another Example: Alignment

- Assign offsets to each variable:

```
int f(int i, int j) {      // i = +8, j = +12
    char x, y, *p;          // x = -1, y = -2, p = -8
    double z;                // z = -16
    char a[5], *s;           // a = -21, s = -28
}
```

- Assume memory accesses should be aligned:
 - A 4-byte object should be aligned on a 4-byte boundary;
 - An 8-byte object should be aligned on an 8-byte boundary;
 - A byte object is always aligned.

Another Example: Alignment

- Assign offsets to each variable:

```
int f(int i, int j) {  
    char x, y, *p;  
    double z;  
    char a[5], *s;  
}
```

