

COEN 175

Lecture 21: Floating Point

Intel Floating Point

- Intel originally used a coprocessor for floating-point.
- The coprocessor uses a stack for computation, so an expression is evaluated as if written in postfix.
- Consider the expression $a * b + c * d$:

```
fildl    a      # load/push a  
fildl    b      # load/push b  
fmulp          # multiply and pop  
fildl    c      # load/push c  
fildl    d      # load/push d  
fmulp          # multiply and pop  
faddp          # add and pop
```

- Any return value is left on the top of the stack.

Modern Floating Point

- Modern Intel processors (since the Pentium III) use SSE (Streaming SIMD Extensions) for floating point.
 - The SSE register names are %xmm0 – %xmm7.
 - They are all **caller-saved** registers.
- The arithmetic opcodes are:
 - `movsd` – move scalar, double-precision
 - `addsd` – add scalars, double-precision
 - `subsd` – subtract scalars, double-precision
 - `mulsd` – multiply scalars, double-precision
 - `divsd` – divide scalars, double-precision
 - `ucomisd` – (unordered) compare scalars, double-precision

Example: Addition

- Assume the registers are %xmm0, %xmm1, etc.
- Assume that all variables are 64-bit reals and are global variables so they can be referred to by name.
- Generate code for $a * b + c * d$.

```
    movsd  a, %xmm0          # load: %xmm0 allocated
    mulsd  b, %xmm0
    movsd  c, %xmm1          # load: %xmm1 allocated
    mulsd  d, %xmm1
    addsd  %xmm1, %xmm0      # %xmm1 deallocated
```

- Compare this with our code for integer arithmetic.

Putting It All Together

- Let's write a complete function for addition.

```
void Add::generate()
{
    // Generate code for both the left child
    // and the right child

    // If the left child is not in a register,
    // then allocate a register and load it

    // Perform the operation such as
    // "addl right, left"

    // If the right operand is in a register,
    // then deallocate it
}
```

Putting It All Together

- Let's write a complete function for addition.

```
void Add::generate()
{
    _left->generate();
    _right->generate();

    if (_left->register == nullptr)
        load(_left, FP(_left) ? fp_getreg() : getreg());

    cout << "\tadd" << suffix(_left);
    cout << _right << ", " << _left << endl;

    assign(_right, nullptr);
    assign(this, _left->register);
}
```

Two sets of registers

Utility Functions

- Let's write the utility functions we've used.

```
# define FP(expr)      ((expr)->type().isReal())
# define BYTE(expr)    ((expr)->type().size() == 1)

static string suffix(Expression *expr) {
    return FP(expr) ? "sd\t" : (BYTE(expr) ? "b\t" : "l\t");
}

Register *fp_getreg() {
    for (unsigned i = 0; i < fp_registers.size(); i++)
        if (fp_registers[i]->_node == nullptr)
            return fp_registers[i];

    load(nullptr, fp_registers[0]);
    return fp_registers[0];
}
```

Floating-Point Comparison

- The `ucomisd` instruction sets the flags register based on the floating-point status word.
- The zero and carry flag are set, but not the sign flag.
- Therefore, we must use the **unsigned** set and jump instructions:
 - `setb/jb` – set/jump if below
 - `seta/ja` – set/jump if above
 - `setbe/jbe` – set/jump if below or equal
 - `setae/jae` – set/jump if above or equal

Example: Comparison

- Assume all variables are 64-bit reals.
- Generate floating-point code for $x > y * z$.

```
movsd      y, %xmm0      # %xmm0 allocated
mulsd      z, %xmm0
movsd      x, %xmm1      # %xmm1 allocated
ucomisd    %xmm0, %xmm1  # %xmm0 and %xmm1 deallocated
seta       %al           # %eax allocated
movzbl    %al, %eax
```

- Note that the result has type `int` and is stored in an integer register, not a floating-point register.

Floating-Point Conversion

- To convert a 32-bit integer to a 64-bit real:
 - `cvtsi2sd` – convert scalar integer to scalar double
- To convert a byte to a 64-bit real, we must first sign extend the byte into a 32-bit integer register.
- To convert a 64-bit real to a 32-bit integer:
 - `cvttsd2si` – convert scalar double to scalar integer
- To convert a 64-bit real to a byte, we first convert the real into a 32-bit integer register and then use the corresponding byte register.

Example: Conversion

- Assume n is a 32-bit integer and x is a 64-bit real.
- Generate code for $n = (\text{char}) (x + n)$.
 - With no coercions: $n = (\text{int}) (\text{char}) (x + (\text{double}) n)$.

```
cvtsi2sd    n, %xmm0
movsd       x, %xmm1
addsd       %xmm0, %xmm1
cvttsd2si   %xmm1, %eax
movsb1      %al, %eax
movl        %eax, n
```

- Note that the cast to a `char` generates no code beyond the conversion to `int`, but we do use `%al` for the next operation.

Floating-Point Literals

- Floating-point literals must be stored in memory and referenced using an assembler label.
- The `.double` directive is used for this purpose:

```
.L0:    .double    3.14159  
.L1:    .double    2.989792e+8
```

- Note that like the `.asciz` directive for strings, the `.double` directive must be used in the `.data` section of the assembly file.

Floating-Point Zero

- A floating-point value of zero is useful for:
 - Determining a truth value;
 - Negating a floating-point value.
- Although floating-point literals must be stored in memory, there is a simple way to generate zero:

```
pxor    %xmm0, %xmm0
```

- Generate code for $-x + y$.

```
pxor    %xmm0, %xmm0
subsd  x, %xmm0
addsd  y, %xmm0
```

Floating-Point Arguments

- Floating-point arguments are pushed on the stack.
 - However, we cannot simply “push” them because they are 8-byte values, not 4-byte values.
 - Instead, we must first adjust the stack pointer and move them ourselves, effectively mimicking a push instruction.
- Generate code for $f(x + y, z)$.

```
    movsd  z, %xmm0
    subl   $8, %esp
    movsd  %xmm0, (%esp)
    movsd  x, %xmm0
    addsd  y, %xmm0
    subl   $8, %esp
    movsd  %xmm0, (%esp)
    call   f
    addl   $16, %esp
```

Floating-Point Return

- For compatibility, a floating-point return value is still placed on the top of the coprocessor stack.
- Since there is no direct path between the SSE registers and the coprocessor, we must use memory.

return $x + y$

```
movsd  x, %xmm0
addsd  y, %xmm0
movsd  %xmm0, t0
fldl   t0
jmp    f.exit
```

$x * f() + y$

```
call   f
fstpl  t1
movsd  x, %xmm0
mulsd  t1, %xmm0
addsd  y, %xmm0
```

Spill and
reload

Spill

Summary

- Within our framework, floating-point values are treated no differently than integer values.
 - Our `load()` function is intelligent enough to handle either integer or floating-point loads.
 - We just need to tell it the correct register.
 - The arithmetic and comparison code generation functions require little to no modification.
- Code generation for type conversion is tricky.
- Function calls and returns require using the coprocessor stack and reloading from memory.