

An Introduction to R

A Programming Environment for Data Analysis and Graphics

Giovanni Seni

September 19, 2018

Table of Contents

The R environment and language

Data Processing

Graphical Procedures

Statistical Models

R and Hadoop

Getting Help

R is an environment for...

Data Manipulation

- ▶ Connecting to data sources
- ▶ Slicing and dicing data

R is an environment for...

Data Manipulation

- ▶ Connecting to data sources
- ▶ Slicing and dicing data

Modeling and Computation

- ▶ Statistical modeling
- ▶ Numerical simulation

R is an environment for...

Data Manipulation

- ▶ Connecting to data sources
- ▶ Slicing and dicing data

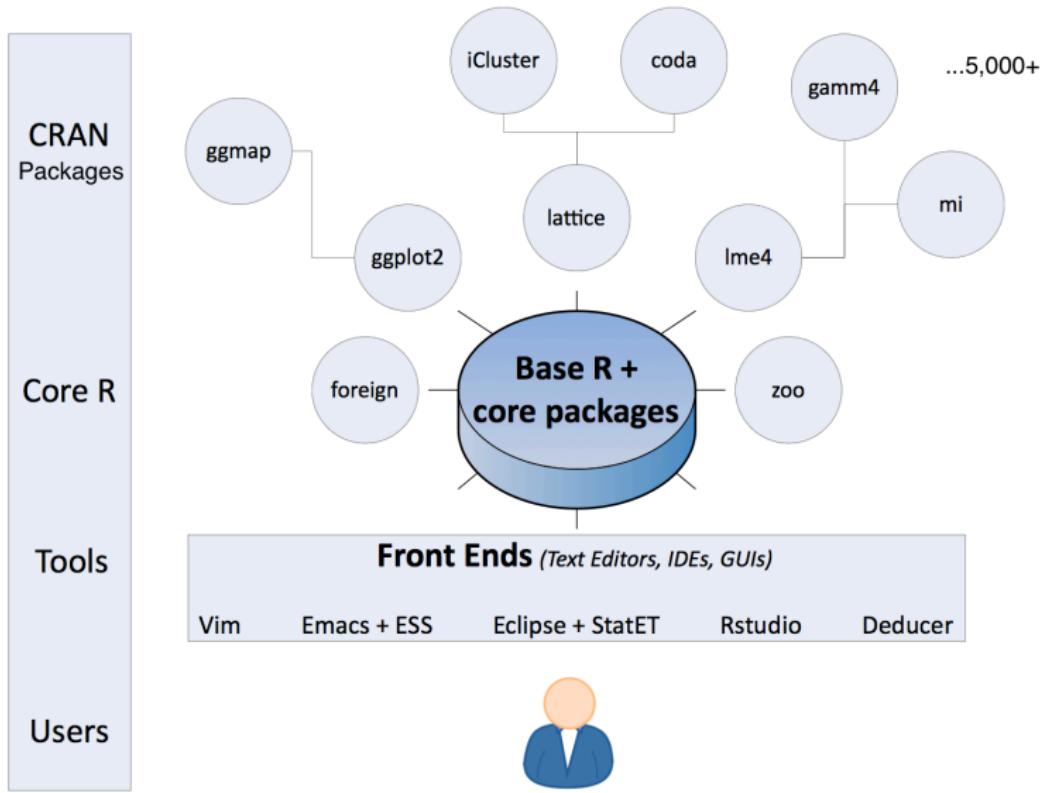
Modeling and Computation

- ▶ Statistical modeling
- ▶ Numerical simulation

Data Visualization

- ▶ Visualizing fit of models
- ▶ Composing statistical graphics

The R Ecosystem



A Sample Session

```
open -a RStudio Diamonds.R
```

The screenshot shows the RStudio interface with the following components:

- Top Bar:** Contains the RStudio logo, a search bar labeled "Go to file/function", and a "Project: (None)" dropdown.
- Left Panel:** A code editor window titled "Diamonds.R x" containing R code. The code reads a CSV file, creates a summary, attaches the data frame, plots Carat Weight vs Price, and creates a bwplot of Price vs Color.

```
1 x <- read.csv("Diamonds.csv")
2 dim(x)
3 colnames(x)
4 summary(x$Price)
5
6 attach(x)
7 plot(CaratWeight, Price, main="",
8       xlab="Carat Weight ", ylab="Price ", pch=19)
9
10 library(lattice)
11 bwplot(Price ~ Color, data = x, xlab = "Color")
12
```
- Environment Tab:** Shows the global environment with a data frame "x" containing 3000 observations.
- Data Tab:** Shows the data frame "x" with 3000 observations.
- Files Tab:** Shows the project structure: Doc > R > R-intro-slides > MyData. It lists three files: Diamonds.R (228 B, Mar 6, 2014, 8:38 AM), Diamonds.csv (116.5 KB, Mar 5, 2014, 5:46 PM), and .Rhistory (30 B, Mar 6, 2014, 8:39 AM).
- Console:** Displays the R session output, including the R prompt (>), the first few lines of the R code, and the message "Type 'q()' to quit R."
- Bottom Right:** Includes standard OS X window control buttons (minimize, maximize, close) and a magnifying glass icon for search.

Simple manipulations; numbers and vectors

- ▶ Simple math

```
> 2+2
```

```
[1] 4
```

- ▶ Storing results in variables

```
> x <- 2+2  ## R syntax for '=' or assignment  
> x^2
```

```
[1] 16
```

Simple manipulations; numbers and vectors

- ▶ Simple math

```
> 2+2
```

```
[1] 4
```

- ▶ Storing results in variables

```
> x <- 2+2  ## R syntax for '=' or assignment  
> x^2
```

```
[1] 16
```

- ▶ Vectorized math

```
> weight <- c(110, 180, 240)  ## vector of 3 weights  
> height <- c(5.5, 6.1, 6.2)  ## vector of 3 heights  
> bmi <- (weight*4.88)/height^2  ## element-wise op
```

Sequences

- ▶ : operator

```
> -5:5
```

```
[1] -5 -4 -3 -2 -1 0 1 2 3 4 5
```

Sequences

- ▶ : operator

```
> -5:5
```

```
[1] -5 -4 -3 -2 -1 0 1 2 3 4 5
```

- ▶ function *seq()*

```
> seq(-1, 1, by=.5)
```

```
[1] -1.0 -0.5 0.0 0.5 1.0
```

Sequences

- ▶ : operator

```
> -5:5
```

```
[1] -5 -4 -3 -2 -1 0 1 2 3 4 5
```

- ▶ function *seq()*

```
> seq(-1, 1, by=.5)
```

```
[1] -1.0 -0.5 0.0 0.5 1.0
```

- ▶ function *rep()*

```
> rep(c(-1, 1), times=2)
```

```
[1] -1 1 -1 1
```

Logical Vectors and Missing Values

- ▶ conditions

```
> heavy <- weight > 200
```

```
[1] FALSE FALSE TRUE
```

Logical Vectors and Missing Values

- ▶ conditions

```
> heavy <- weight > 200
```

```
[1] FALSE FALSE TRUE
```

- ▶ NA : “not available”

```
> z <- c(1:3,NA)
```

```
[1] 1 2 3 NA
```

Logical Vectors and Missing Values

- ▶ conditions

```
> heavy <- weight > 200
```

```
[1] FALSE FALSE TRUE
```

- ▶ NA : “not available”

```
> z <- c(1:3,NA)
```

```
[1] 1 2 3 NA
```

- ▶ function *is.na(x)*

```
> ind <- is.na(z)
```

```
[1] FALSE FALSE FALSE TRUE
```

Character vectors

Character strings are entered using either matching double ("") or single ('') quotes, but are printed using double quotes

```
> labs <- paste(c('X', 'Y'), 1:3, sep="-")  
[1] "X-1" "Y-2" "X-3"
```

Note that recycling of short lists takes place...

Index Vectors

```
> x[1:10]
```

selects the first 10 elements of x (assuming $\text{length}(x)$ is not less than 10)

Index Vectors

```
> x[1:10]
```

selects the first 10 elements of x (assuming $\text{length}(x)$ is not less than 10)

```
> x[-(1:5)]
```

gives all but the first five elements of x

Index Vectors

```
> x[1:10]
```

selects the first 10 elements of x (assuming length(x) is not less than 10)

```
> x[-(1:5)]
```

gives all but the first five elements of x

Objects can have a names attribute to identify its components

```
> fruit <- c(5, 10, 1, 20)
```

```
> names(fruit) <- c("orange", "banana", "apple", "peach")
```

```
> fruit[c("apple", "orange")]
```

apple orange 1 5

Factors, or categorical variables

Factors in R are stored as a vector of integer values with a corresponding set of character values to use when the factor is displayed.

- ▶ Create a factor corresponding to weight, with three equally spaced levels:

```
> weight.factor <- cut(weight, 3)
> summary(weight.factor)

## (115,131] (131,148] (148,164]
##           6           5           4
```

Factors, or categorical variables

Factors in R are stored as a vector of integer values with a corresponding set of character values to use when the factor is displayed.

- ▶ Create a factor corresponding to weight, with three equally spaced levels:

```
> weight.factor <- cut(weight, 3)
> summary(weight.factor)

## (115,131] (131,148] (148,164]
##          6           5           4
```

- ▶ The *class* and *levels* of a factor

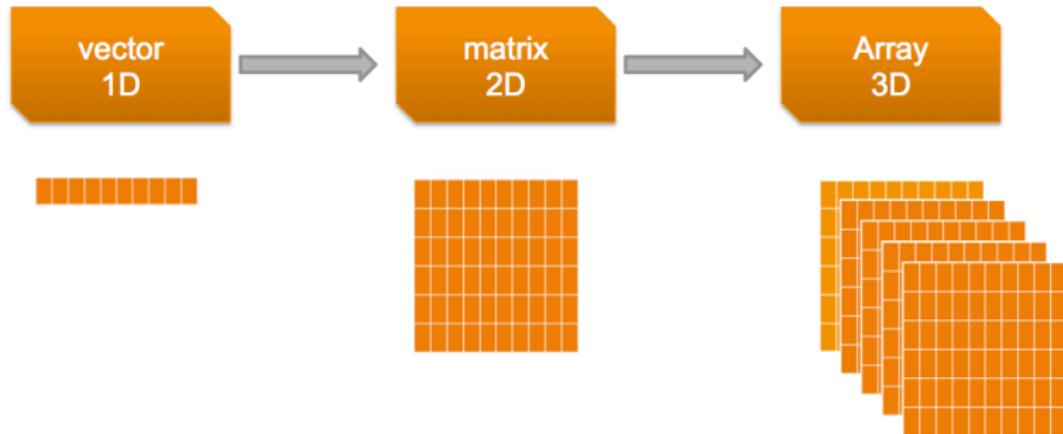
```
> class(weight.factor)
```

```
[1] "factor"
```

```
> levels(weight.factor)
```

```
[1] "(115,131]" "(131,148]" "(148,164]"
```

Arrays and matrices



Arrays and matrices

► The *matrix()* function

```
> M <- matrix(1:9, 3, 3, byrow = T)

##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
## [3,]    7    8    9
```

Arrays and matrices

- ▶ The *matrix()* function

```
> M <- matrix(1:9, 3, 3, byrow = T)

##      [,1] [,2] [,3]
## [1,]     1     2     3
## [2,]     4     5     6
## [3,]     7     8     9
```

- ▶ Negative subscripts remove elements:

```
> M[-1, -2]

##      [,1] [,2]
## [1,]     4     6
## [2,]     7     9
```

Arrays and matrices

Matrix Facilities

- ▶ Row and column name assignments:

```
> colnames(M) <- c("C1", "C2")
> rownames(M) <- c("R1", "R2")

##      C1  C2
## R1    4   6
## R2    7   9
```

Arrays and matrices

Matrix Facilities

- ▶ Row and column name assignments:

```
> colnames(M) <- c("C1", "C2")
> rownames(M) <- c("R1", "R2")

##      C1  C2
## R1    4   6
## R2    7   9
```

- ▶ Subsetting with names:

```
> M[, "C1"]

## R1  R2
##  4   7
```

Arrays and matrices

Matrix Facilities

- ▶ Matrix dimensions:

```
> nrow(A)  
> ncol(A)
```

Arrays and matrices

Matrix Facilities

- ▶ Matrix dimensions:

```
> nrow(A)  
> ncol(A)
```

- ▶ Matrix product:

```
> A %*% B
```

Arrays and matrices

Matrix Facilities

- ▶ Matrix dimensions:

```
> nrow(A)  
> ncol(A)
```

- ▶ Matrix product:

```
> A %*% B
```

- ▶ Matrix transpose:

```
> t(A)
```

Arrays and matrices

Matrix Facilities

- ▶ Matrix dimensions:

```
> nrow(A)  
> ncol(A)
```

- ▶ Matrix product:

```
> A %*% B
```

- ▶ Matrix transpose:

```
> t(A)
```

- ▶ Matrix inverse:

```
> solve(A)
```

Arrays and matrices

Matrix Facilities

- ▶ `cbind()`: forms matrices by binding together vectors horizontally, or column-wise

```
> X1 <- c(1, 2)
> X2 <- c(3, 4)
> X <- cbind(0, X1, X2)

##          X1 X2
## [1,] 0  1  3
## [2,] 0  2  4
```

Arrays and matrices

Matrix Facilities

- ▶ *cbind()*: forms matrices by binding together vectors horizontally, or column-wise

```
> X1 <- c(1, 2)
> X2 <- c(3, 4)
> X <- cbind(0, X1, X2)

##          X1 X2
## [1,] 0  1  3
## [2,] 0  2  4
```

- ▶ *rbind()*: corresponding operation for rows

Lists

An R *list* is an object consisting of an ordered collection of objects of possibly different types

```
> Lst <- list(name="Fred", wife="Mary",
+                 no.children=3,
+                 child.ages=c(4,7,9))
```

Lists

An R *list* is an object consisting of an ordered collection of objects of possibly different types

```
> Lst <- list(name="Fred", wife="Mary",
+                 no.children=3,
+                 child.ages=c(4,7,9))
```

- ▶ Accessing components:

```
> Lst$name == Lst[[1]]
```

```
[1] TRUE
```

```
> Lst$child.ages[1] == Lst[[4]][1]
```

```
[1] TRUE
```

Data Frames

A *data frame* is a list where components are vectors, factors, lists, or other data frames

```
> char <- letters[1:5]
> m <- 1:5
> log <- c(TRUE, FALSE, TRUE, FALSE, TRUE)
> df <- data.frame(char, m, log)

##   char m   log
## 1   a 1  TRUE
## 2   b 2 FALSE
## 3   c 3  TRUE
## 4   d 4 FALSE
## 5   e 5  TRUE
```

Data Frames

A *data frame* is a list where components are vectors, factors, lists, or other data frames

```
> char <- letters[1:5]
> m <- 1:5
> log <- c(TRUE, FALSE, TRUE, FALSE, TRUE)
> df <- data.frame(char, m, log)

##   char m   log
## 1   a 1  TRUE
## 2   b 2 FALSE
## 3   c 3  TRUE
## 4   d 4 FALSE
## 5   e 5  TRUE
```

A data frame may for many purposes be regarded as a matrix

Getting Data In/Out

- ▶ from Files

```
> x <- read.csv("Diamonds.csv")
> x <- read.table("Diamonds.txt", header=T, sep=", ")
> x <- read.xlsx(Diamonds.xlsx, 2) # read the second sheet
```

Getting Data In/Out

- ▶ from Files

```
> x <- read.csv("Diamonds.csv")
> x <- read.table("Diamonds.txt", header=T, sep=",")
> x <- read.xlsx(Diamonds.xlsx, 2) # read the second sheet
```

- ▶ from Databases

```
> library(RODBC)
> con <- odbcConnect(db.dsn, uid="", pwd="")
> x <- sqlQuery(con, "SELECT * FROM Diamonds")
> close(con)
```

Getting Data In/Out

- ▶ from Files

```
> x <- read.csv("Diamonds.csv")
> x <- read.table("Diamonds.txt", header=T, sep=",")
> x <- read.xlsx(Diamonds.xlsx, 2) # read the second sheet
```

- ▶ from Databases

```
> library(RODBC)
> con <- odbcConnect(db.dsn, uid="", pwd="")
> x <- sqlQuery(con, "SELECT * FROM Diamonds")
> close(con)
```

- ▶ from R objects

```
> load("Diamonds.RData")
```

Getting Data In/Out

- ▶ from Files

```
> x <- read.csv("Diamonds.csv")
> x <- read.table("Diamonds.txt", header=T, sep=", ")
> x <- read.xlsx(Diamonds.xlsx, 2) # read the second sheet
```

- ▶ from Databases

```
> library(RODBC)
> con <- odbcConnect(db.dsn, uid="", pwd="")
> x <- sqlQuery(con, "SELECT * FROM Diamonds")
> close(con)
```

- ▶ from R objects

```
> load("Diamonds.RData")
```

- ▶ from the Web

```
> con <- url("http://... ")
> x <- read.csv(con)
```

Control Flow Basics

- ▶ *for loops, repeat and while*

```
> for (i in 1:nrow(x)) {  
+   print(x[i,])  
+ }
```

Warning: `for()` loops are used in R code much less often than in compiled languages. Often *apply*-type functions preferred.

Control Flow Basics

- ▶ *for loops, repeat and while*

```
> for (i in 1:nrow(x)) {  
+   print(x[i,])  
+ }
```

Warning: `for()` loops are used in R code much less often than in compiled languages. Often *apply*-type functions preferred.

- ▶ *if statements*

```
> if (expr_1) expr_2 else expr_3  
> ifelse(condition, a, b)
```

operators `&&` and `||` are often used as part of the condition

Control Flow Basics

- ▶ *for loops, repeat and while*

```
> for (i in 1:nrow(x)) {  
+   print(x[i,])  
+ }
```

Warning: `for()` loops are used in R code much less often than in compiled languages. Often *apply*-type functions preferred.

- ▶ *if statements*

```
> if (expr_1) expr_2 else expr_3  
> ifelse(condition, a, b)
```

operators `&&` and `||` are often used as part of the condition

- ▶ *break and next statements available*

Functions

A function is defined by an assignment of the form

```
> name <- function(arg_1,arg_2, ...) {  
+   expression(s)  
+ }
```

Functions

A function is defined by an assignment of the form

```
> name <- function(arg_1,arg_2, ...) {  
+   expression(s)  
+ }
```

if there is a function fun1 defined by

```
> fun1 <- function(data, data.frame, graph, limit) {...}
```

then the function may be invoked in several ways, for example

```
> ans <- fun1(d, df, TRUE, 20)  
> ans <- fun1(d, df, graph=TRUE, limit=20)  
> ans <- fun1(data=d, limit=20, graph=TRUE,  
+               data.frame=df)
```

Navigating within the R Environment

- ▶ Listing all variables

```
> ls()
```

Navigating within the R Environment

- ▶ Listing all variables

```
> ls()
```

- ▶ Examining a variable x

```
> str(x)
> class(x)
> typeof(x)
> attributes(x)
> head(x)
> tail(x)
```

Navigating within the R Environment

- ▶ Listing all variables

```
> ls()
```

- ▶ Examining a variable x

```
> str(x)
> class(x)
> typeof(x)
> attributes(x)
> head(x)
> tail(x)
```

- ▶ Removing variables

```
> rm(x)
> rm(list=ls())
```

Navigating within the R Environment

- ▶ Listing all variables

```
> ls()
```

- ▶ Examining a variable x

```
> str(x)
> class(x)
> typeof(x)
> attributes(x)
> head(x)
> tail(x)
```

- ▶ Removing variables

```
> rm(x)
> rm(list=ls())
```

- ▶ Get or set working directory

```
> getwd()
> setwd(dir)
```

Data Processing

Data Subsetting

- ▶ **which**: used to identify values in a vector or array that satisfy a list of criteria
e.g. which car in *mtcars* has the highest horse power:

```
> which(mtcars$hp == max(mtcars$hp))
```

```
[1] 31
```

Data Processing

Data Subsetting

- ▶ *which*: used to identify values in a vector or array that satisfy a list of criteria

e.g. which car in *mtcars* has the highest horse power:

```
> which(mtcars$hp == max(mtcars$hp))
```

```
[1] 31
```

- ▶ *subset*: allows us to do the same thing, but different notation

```
> subset(mtcars, mtcars$mpg > 20 & mtcars$cyl == 4,  
+         select = c("mpg", "cyl", "disp"))
```

Data Processing

Data Aggregation

- ▶ `table`: build a contingency table of the counts at each combination of factor levels.

```
>  
table(mtcars[, "gear"])  
##  
##   3   4   5  
## 15 12  5
```

```
> table(mtcars[, c("gear", "cyl")])  
##          cyl  
##    gear  4   6   8  
##        3   1   2 12  
##        4   8   4   0  
##        5   2   1   2
```

Data Processing

Data Aggregation

- ▶ `table`: build a contingency table of the counts at each combination of factor levels.

```
>  
table(mtcars[, "gear"])  
##  
##   3   4   5  
## 15 12  5
```

```
> table(mtcars[, c("gear", "cyl")])  
##          cyl  
##          gear  4   6   8  
##            3   1   2 12  
##            4   8   4   0  
##            5   2   1   2
```

- ▶ `xtabs` allows you to do the same thing, but using a formula

```
> xtabs(~gear+cyl, mtcars)
```

Data Processing

Data Aggregation

- ▶ `aggregate`: Splits the data into subsets, computes summary statistics for each (like a sql group by)

```
> aggregate(hp~cyl+gear, data=mtcars, mean)

##      cyl gear      hp
## 1      4    3 97.0000
## 2      6    3 107.5000
## 3      8    3 194.1667
## 4      4    4  76.0000
## 5      6    4 116.5000
## 6      4    5 102.0000
## 7      6    5 175.0000
## 8      8    5 299.5000
```

Data Processing

Data Aggregation

- ▶ *reshape*: from repeated measurements in separate records to repeated measurements in separate columns

```
> df.wide <- reshape(movie.data, idvar="user.id",
+                      timevar="item.id",
+                      direction="wide")
```

| user.id | item.id | rating |
|---------|---------|--------|
| 196 | 242 | 3 |
| 196 | 302 | 1 |
| 186 | 302 | 3 |
| ... | ... | ... |

| row.names | user.id | rating.242 | rating.302 |
|-----------|---------|------------|------------|
| 1 | 196 | 3 | 1 |
| 2 | 186 | NA | 3 |
| ... | ... | ... | ... |

Data Processing

Merging

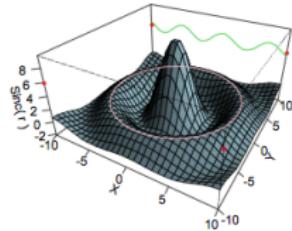
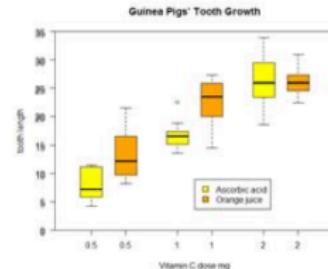
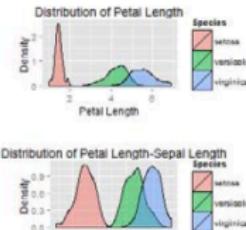
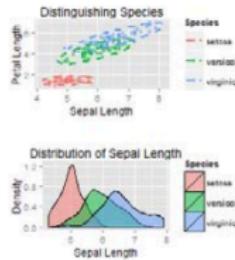
merge: equivalent to database join operations to combine data sets

```
> df.all <- merge(df.1, df.2, by = "id", all = TRUE)
```

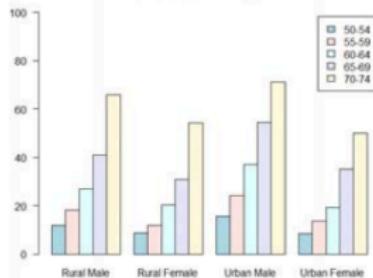
- ▶ `all = FALSE`: gives a *natural* (inner) join
- ▶ `all.x = TRUE`: gives a left (outer) join
- ▶ `all.y = TRUE`: gives a right (outer) join
- ▶ `all = TRUE`: gives a (full) outer join

Graphical Procedures

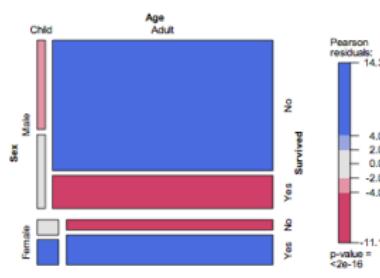
R does graphs



Death Rates in Virginia

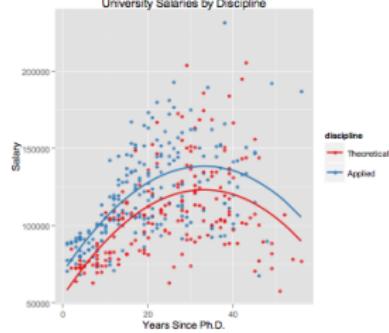


Survival on the Titanic



Pearson residuals: 14.3
p-value = 2e-16

University Salaries by Discipline



Graphical Procedures

Basic Graphs

plot() is a generic function: the type of plot produced is dependent on the type or class of the first argument

```
> plot(CaratWeight, Price,  
+       main="",  
+       xlab="Carat Weight",  
+       ylab="Price ", pch=19)
```

Graphical Procedures

Basic Graphs

`plot()` is a generic function: the type of plot produced is dependent on the type or class of the first argument

```
> plot(CaratWeight, Price,  
+       main="",  
+       xlab="Carat Weight",  
+       ylab="Price ", pch=19)
```

```
> plot(Color)
```

generates a bar plot of diamond's *Color*

Graphical Procedures

Basic Graphs

`plot()` is a generic function: the type of plot produced is dependent on the type or class of the first argument

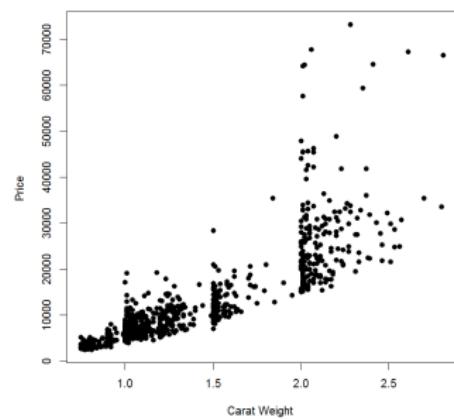
```
> plot(CaratWeight, Price,  
+       main="",  
+       xlab="Carat Weight",  
+       ylab="Price ", pch=19)
```

```
> plot(Color)
```

generates a bar plot of diamond's *Color*

```
> plot(Color, Price)
```

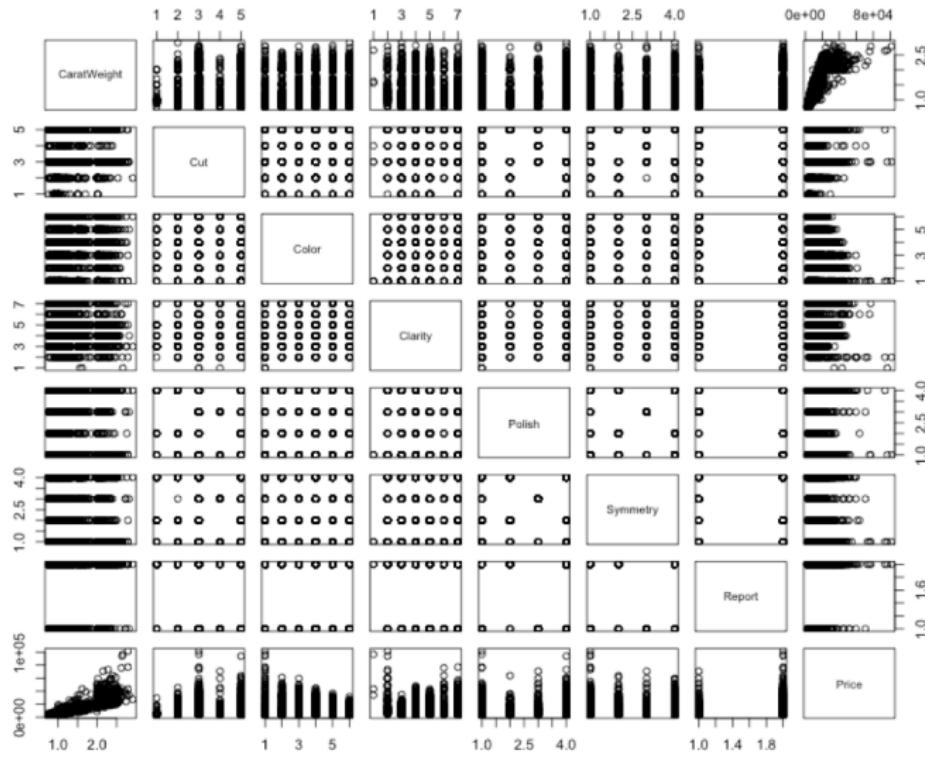
produces boxplots of *Price* for each level of *Color*



Graphical Procedures

Scatter Plot Matrix

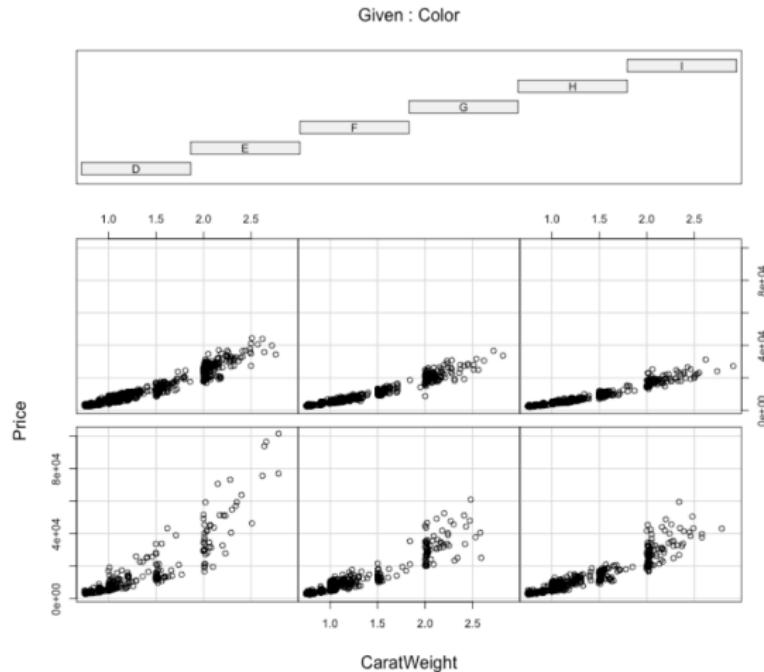
```
> pairs(x[, 2:9])
```



Graphical Procedures

Conditioning Plots

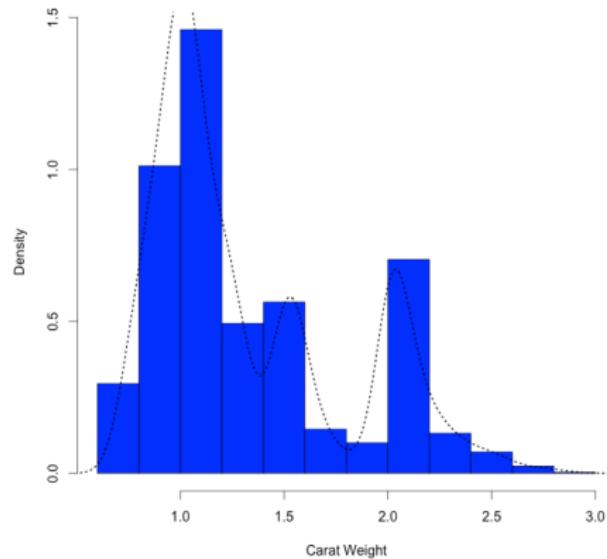
```
> coplot(Price ~ CaratWeight | Color)
```



Graphical Procedures

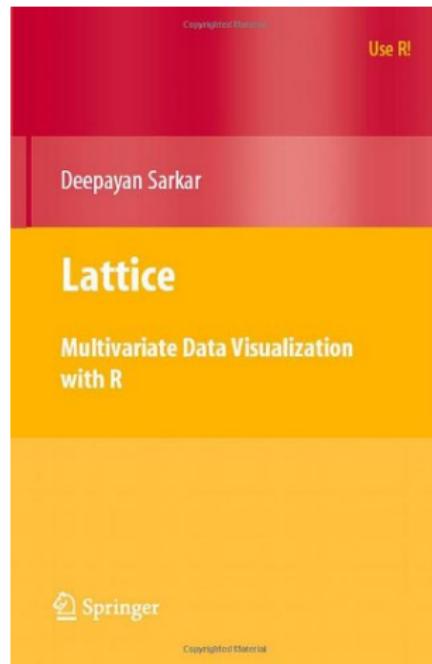
Histograms and Density Plots

```
> hist(x$CaratWeight, breaks=12, col="blue", main="",
+       xlab="Carat Weight", prob=T)
> lines(density(x$CaratWeight), lty="dotted", lwd=2)
```



Graphical Procedures

Lattice vs ggplot



Graphical Procedures

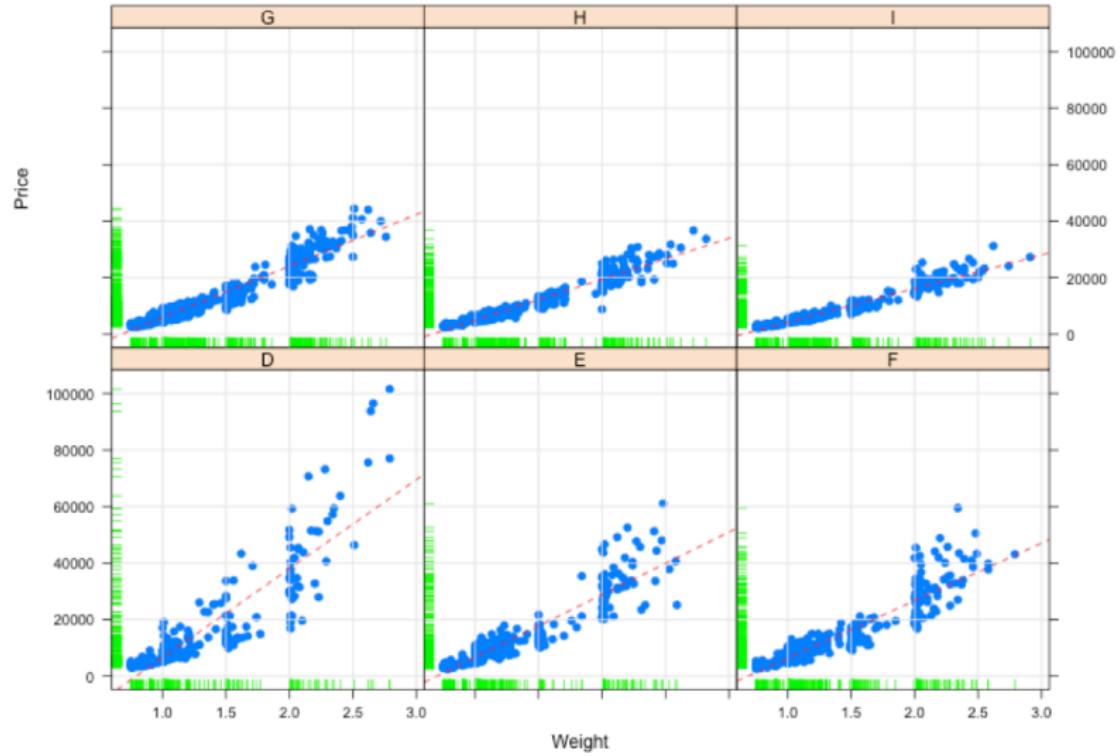
Customizing Lattice Graphs

```
> library(lattice)
> xyplot(Price ~ CaratWeight | Color, data = x,
+         layout=c(3,3),
+         xlab="Weight", ylab="Price",
+         panel=mypanel)
```

```
> mypanel <- function(x, y) {
+   panel.xyplot(x, y, pch=19)
+   panel.rug(x, y, col="green")
+   panel.grid(h=-1, v=-1)
+   panel.lmline(x, y, col="red", lwd=1, lty=2)
+ }
```

Graphical Procedures

Customizing Lattice Graphs



Graphical Procedures

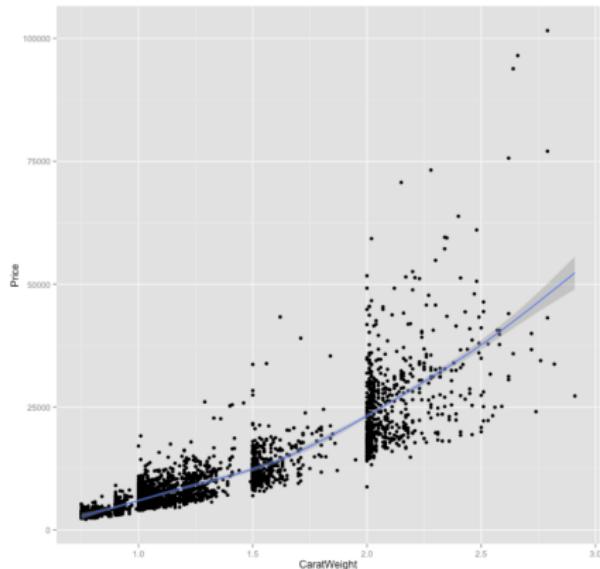
Lattice Graphs

| Graph.type | Function | Formula.examples |
|------------------------------|---------------|--------------------------|
| 1 3D contour plot | contourplot() | $z \sim x * y$ |
| 2 3D level plot | levelplot() | $z \sim y * x$ |
| 3 3D scatter plot | cloud() | $z \sim x * y A$ |
| 4 3D wireframe graph | wireframe() | $z \sim y * x$ |
| 5 Bar chart | barchart() | $x \sim A$ or $A \sim x$ |
| 6 Box plot | bwplot() | $x \sim A$ or $A \sim x$ |
| 7 Dot plot | dotplot() | $\sim x A$ |
| 8 Histogram | histogram() | $\sim x$ |
| 9 Kernel density plot | densityplot() | $\sim x A * B$ |
| 10 Parallel coordinates plot | parallel() | dataframe |
| 11 Scatter plot | xyplot() | $y \sim x A$ |
| 12 Scatter plot matrix | splom() | dataframe |
| 13 Strip plots | stripplot() | $A \sim x$ or $x \sim A$ |

Graphical Procedures

GGplot scatter plot with loess fit

```
> library(ggplot2)
> qplot(Price, CaratWeight, data = x,
+       geom = c("point", "smooth"))
```



Statistical Models

Fit a Linear Model

```
> diamonds.lm <- lm(Price ~ ., data = diamonds.df)
```

- ▶ examine it

```
> summary(diamonds.lm)
```

Statistical Models

Fit a Linear Model

```
> diamonds.lm <- lm(Price ~ ., data = diamonds.df)
```

- ▶ examine it

```
> summary(diamonds.lm)
```

- ▶ plot it

```
> plot(diamonds.lm)
```

```
> plot(density(resid(diamonds.lm)))
```

Statistical Models

Fit a Linear Model

```
> diamonds.lm <- lm(Price ~ ., data = diamonds.df)
```

- ▶ examine it

```
> summary(diamonds.lm)
```

- ▶ plot it

```
> plot(diamonds.lm)
```

```
> plot(density(resid(diamonds.lm)))
```

- ▶ add interactions

```
> diamonds.lm <- lm(Price ~ CaratWeight + Cut*Color,  
+ data = diamonds.df)
```

Statistical Models

Fit a Linear Model

```
> diamonds.lm <- lm(Price ~ ., data = diamonds.df)
```

- ▶ examine it

```
> summary(diamonds.lm)
```

- ▶ plot it

```
> plot(diamonds.lm)
```

```
> plot(density(resid(diamonds.lm)))
```

- ▶ add interactions

```
> diamonds.lm <- lm(Price ~ CaratWeight + Cut*Color,  
+ data = diamonds.df)
```

- ▶ predict

```
> y.hat <- predict(diamonds.lm, newdata)
```

Statistical Models

Fit a Logistic Regression

- ▶ need a binary outcome variable

```
> diamonds.df$Price01 <- diamonds.df$Price > 10000  
> summary(diamonds.df$Price01)  
  
##      Mode   FALSE    TRUE  
## logical 1837    1163
```

Statistical Models

Fit a Logistic Regression

- ▶ need a binary outcome variable

```
> diamonds.df$Price01 <- diamonds.df$Price > 10000  
> summary(diamonds.df$Price01)  
  
##      Mode   FALSE    TRUE  
## logical 1837    1163
```

- ▶ fit model... exclude original *Price*

```
> diamonds.glm <- glm(Price01 ~., data =  
+                           subset(diamonds.df,  
+                           select = -c(Price, ID)))
```

Statistical Models

Fit a Decision Tree

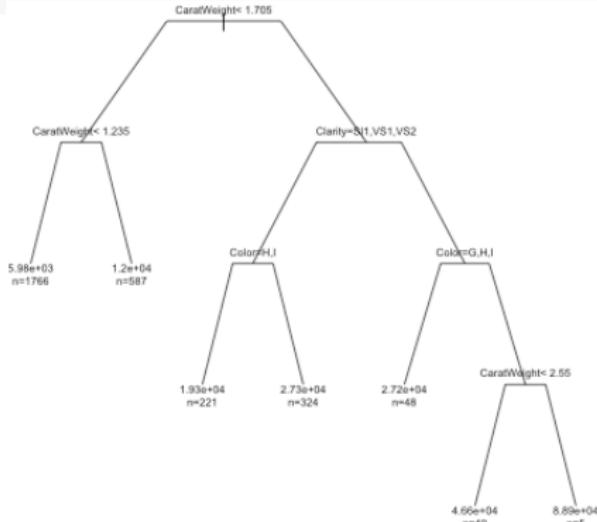
```
> library(rpart)
> set.seed(123)
> diamonds.rpa <- rpart(Price ~ .,
+                         method = "anova",
+                         data = diamonds.df[,-1],
+                         cp = 0,
+                         minsplit=2,
+                         minbucket=1,
+                         maxdepth=7)
```

- ▶ *cp* pre-pruning parameter: Any split that does not decrease the overall lack of fit by a factor of *cp* is discarded

Statistical Models

Fit a Decision Tree

```
> diamond.rpa1 <- prune(diamonds.rpa, cp = 0.02)
> plot(diamond.rpa1, branch = 0.4, uniform = T,
+       compress = T)
> text(diamond.rpa1, use.n = T, digits=3, cex = 0.5,
+       pretty=0)
```



R and Hadoop

- ▶ Install Cloudera ODBC Driver for Hive
- ▶ Get Kerberos authentication ticket

```
kinit -V user-admin-name
```

- ▶ Within R

```
> library(RODBC)  
> ch = odbcConnect("HaaS_Hive_DSN")  
> sqlQuery(ch, "select * from db_name.table_name")
```

Getting Help

| Function | Action |
|--------------------|---|
| help.start() | General help |
| ?foo | Help on function foo |
| ??foo | Search the help system for instances of the string foo |
| example("foo") | Examples of function foo |
| RSiteSearch("foo") | Search for the string foo in online help manuals and archived mailing lists |
| apropos("foo") | List all available functions with foo in their name |
| data() | List all available datasets for currently loaded packages |
| vignette(topic) | Open a vignette for topic or package |