

COEN281 -- Introduction to Pattern Recognition and Data Mining

Lecture 7: Neural Networks

Instructor: Dr. Giovanni Seni
GSeni@scu.edu

*Department of Computer Engineering
Santa Clara University*

Fall/18

Syllabus

| | |
|---------|---|
| Week 1 | Introduction; R (Ch.1) |
| Week 2 | Bayesian Decision Theory (Ch.2; DHS: 2.1-2.6, 2.9) Parameter Estimation (DHS: 3.1-3.4) |
| Week 3 | Linear Discriminant Functions (Ch.3&4; DHS: 3.8.2, 5.1-5.8) Regularization (Ch.6; SE: Ch.3) |
| Week 4 | Neural Networks (DHS: 6.1-6.6, 6.8); Deep Learning |
| Week 5 | Support Vector Machines (Ch.9) |
| Week 6 | Decision Trees (Ch. 8.1; DHS: 8.3; Ch 2 SE) |
| Week 7 | Ensemble Methods (Ch. 8.2; SE: Ch 4, 5) |
| Week 8 | Clustering (Ch. 10; DHS: 10.6, 10.7) Clustering (DHS: 10.9); How many clusters are there? (DHS: 10.10) |
| Week 9 | Non-metric: Association Rules Collaborative Filtering |
| Week 10 | Text Retrieval; Other topics |

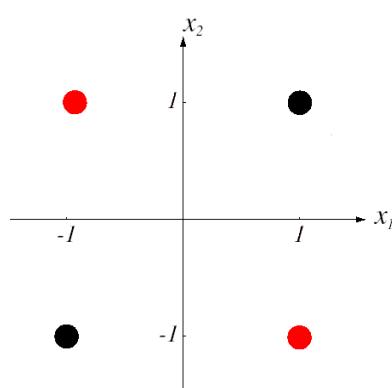
Overview

- Introduction
 - Generalized linear discriminant functions
 - Feedforward neural networks
 - Biological motivation
- Feedforward Networks
 - Net and output activation; activation functions
 - Relationship to Logistic Regression
 - Expressive power
- Network Learning
 - Overview
 - *Backpropagation* algorithm

Introduction

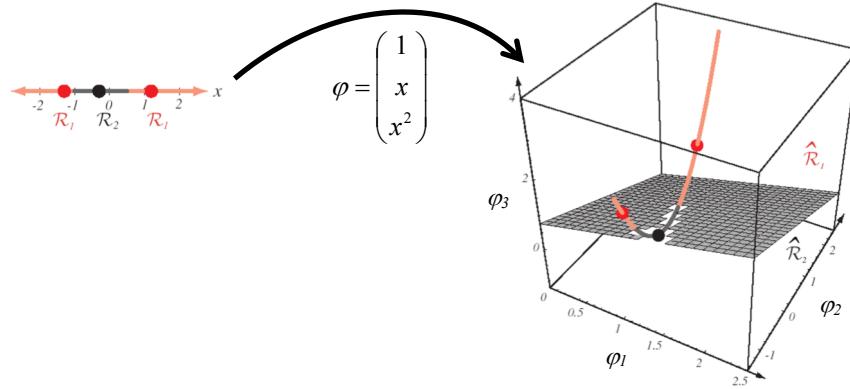
Generalized Linear Discriminant

- There are many problems for which linear discriminants are insufficient
 - XOR



Neural Networks

Generalized Linear Models (2)



- Idea: use functions $\varphi(x)$ to map points in x -space to points in φ -space where a linear function suffices

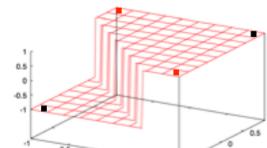
$$g(\mathbf{x}) = \sum_{i=1}^{\hat{d}} w_i \varphi_i(\mathbf{x})$$

Neural Networks

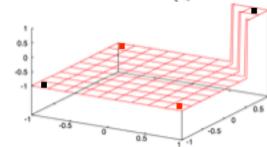
Generalized Linear Models (3)

- NN idea - employ multiple copies of a single nonlinear function $\bar{\varphi}$

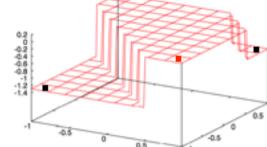
$$y_1 = \varphi_1(\mathbf{x}) = \begin{cases} 1 & \text{if } (x_1 + x_2 + 0.5) \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



$$y_2 = \varphi_2(\mathbf{x}) = \begin{cases} 1 & \text{if } (x_1 + x_2 - 1.5) \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



$$g(\mathbf{x}) = 0.7 \times y_1 - 0.4 \times y_2 - 1$$

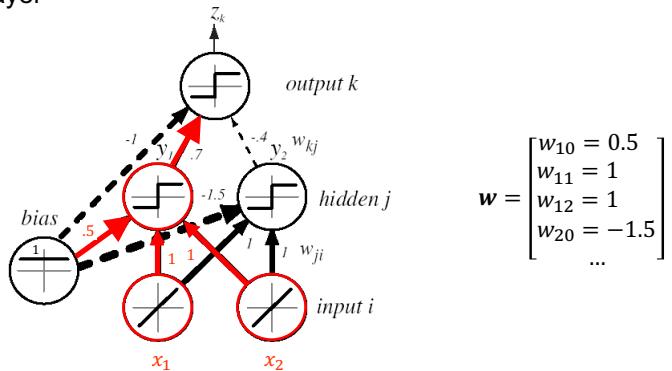


$\bar{\varphi} : \varphi(\mathbf{x}) = \text{sign function}$

Introduction

Feedforward Neural Networks

- Each $\varphi()$ function is associated with a “unit” in a directed graph
 - units are arranged into “layers” – e.g., input layer, hidden layer, and output layer



COEN281

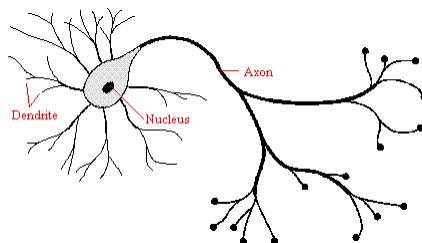
GSeni@scu.edu

9

Neural Networks

Biological Inspiration

- Human brain contains ~ 100 billion neurons, each connected to $\sim 10,000$ other cells



- Cell body (CPU): combines signals
- Dendrite: input bus
- Axon (cable): will transport the activation signal to neurons at different locations

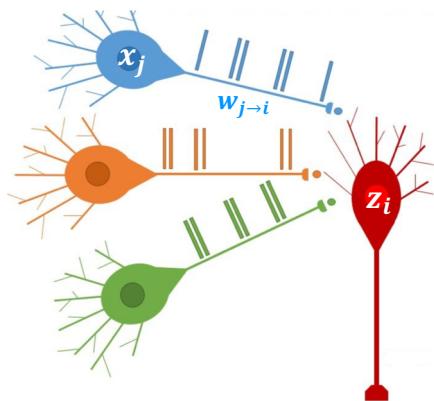
COEN281

GSeni@scu.edu

11

Neural Networks

Simple Computation



$$x_j \xrightarrow{w_{j \rightarrow i}} z_i$$

$$x_j = \begin{cases} 1 & \text{neuron fires} \\ 0 & \text{else} \end{cases}$$

$$w_{j \rightarrow i} = \begin{cases} > 0 & \text{stimuli} \\ < 0 & \text{inhibit} \end{cases}$$

$$z_i = \begin{cases} 1 & \sum w_{j \rightarrow i} x_j > w_{j0} \\ 0 & \text{else} \end{cases}$$

COEN281

GSeni@scu.edu

12

Feedforward Network

Net and Output Activation

- Hidden layer:

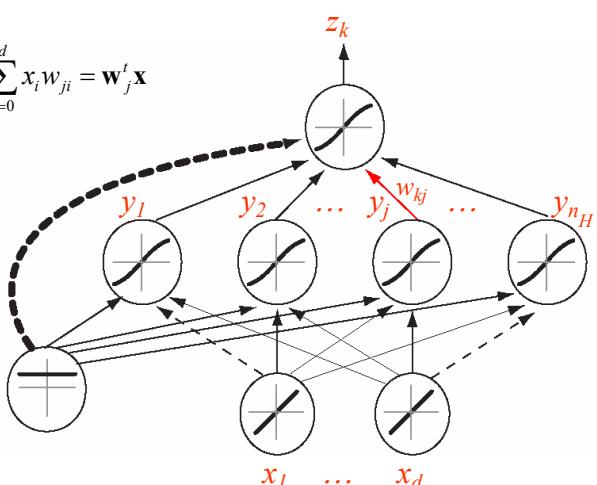
$$net_j = \sum_{i=1}^d x_i w_{ji} + w_{j0} = \sum_{i=0}^d x_i w_{ji} = \mathbf{w}_j^t \mathbf{x}$$

$$y_j = f(net_j)$$

- Output layer:

$$net_k = \sum_{j=0}^{n_H} y_j w_{kj} = \mathbf{w}_k^t \mathbf{y}$$

$$z_k = f(net_k)$$



COEN281

GSeni@scu.edu

13

Feedforward Network

Activation Function $f(\theta) = \varphi(\theta)$

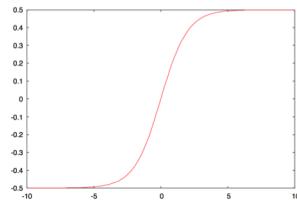
- Required characteristics
 - Nonlinear
 - Continuous and differentiable
- Desirable
 - Saturates
 - Keeps the weights and activations bounded
 - Monotonic
 - Avoids additional and unnecessary local extrema in error surface
 - Linear for small values of net
 - Allows linear model if appropriate

Feedforward Network

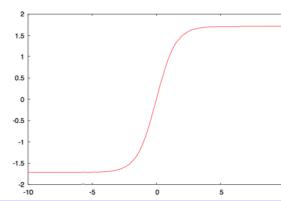
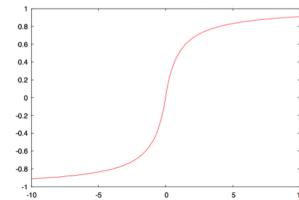
Activation Function (2)

- Centered on zero and anti-symmetric preferred for “faster” learning

$$f(net) = \frac{1}{1 + e^{-net}} - 0.5$$



$$f(net) = \frac{net}{1 + |net|}$$



$$f(net) = 1.716 \times \tanh\left(\frac{2}{3} \times net\right)$$

Feedforward Network

Relationship to Logistic Regression

- Can view NN as a “recursive” logistic regression
 - Each of the predictors in our logistic regression is itself a logistic regression on the data

$$P(\omega_1 | \mathbf{x}) = \phi(\boldsymbol{\beta}' \mathbf{x}) \quad \Rightarrow \quad P(\omega_1 | \mathbf{x}) = \phi(\boldsymbol{\beta}' \boldsymbol{\Theta}(\mathbf{x}))$$

$$\text{where } \boldsymbol{\Theta}_1(\mathbf{x}) = \phi(\boldsymbol{\beta}'_1 \mathbf{x})$$

...

$$\boldsymbol{\Theta}_{n_H}(\mathbf{x}) = \phi(\boldsymbol{\beta}'_{n_H} \mathbf{x})$$

and

$$\phi(z) = \frac{e^z}{1+e^z} = \frac{1}{1+e^{-z}}$$

Feedforward Network

Expressive Power

- Class of functions computed by three-layer NN

$$g_k(x) \equiv z_k = f \left(\sum_{j=1}^{n_H} w_{kj} f \left(\sum_{i=1}^d w_{ji} x_i + w_{j0} \right) + w_{k0} \right)$$

- Can every discriminant function be implemented by a network as described by this equation?

- Yes (due to A. Kolmogorov). Any continuous function $g(x)$ defined on the unit hypercube I^n ($I=[0,1]$ and $n \geq 2$) can be represented in the form:

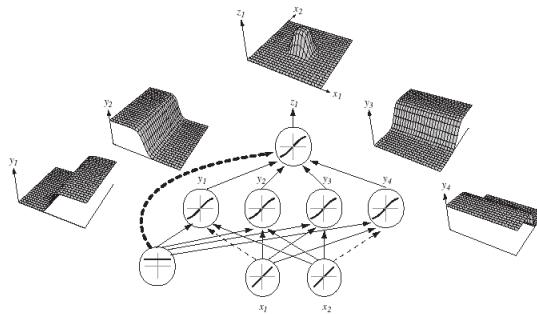
$$g(\mathbf{x}) = \sum_{j=1}^{2n+1} \Xi_j \left(\sum_{i=1}^d \Psi_{ij}(x_i) \right)$$

for properly chosen functions Ξ_j and Ψ_{ij}

Feedforward Network

Expressive Power (2)

- Kolmogorov's Th. is of little practical use
 - Ξ_j and Ψ_{ij} can be extremely complex and not smooth
 - Doesn't tell us how to find the nonlinearities based on the data
- A more intuitive (Fourier-like) proof



Network Learning

Overview

- Goal: set weights w_{kj} and w_{ji} based on the training patterns and the desired outputs
- Supervised training: for every input pattern there is a target vector t
 - E.g., $t_k(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \omega_k \\ 0 & \text{otherwise} \end{cases}$
- Criterion function (on a single pattern)

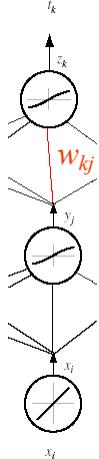
$$J(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^c (t_k - z_k)^2 = \frac{1}{2} \|\mathbf{t} - \mathbf{z}\|^2$$

- Gradient descent step: $\mathbf{w}(k+1) = \mathbf{w}(k) - \eta \nabla J(\mathbf{w})$

Network Learning

Backpropagation

- Hidden-to-output weights w_{kj}



$$\frac{\partial J}{\partial w_{kj}} = \frac{\partial J}{\partial \text{net}_k} \cdot \frac{\partial \text{net}_k}{\partial w_{kj}} = -\delta_k \frac{\partial \text{net}_k}{\partial w_{kj}}$$

$$\delta_k = -\frac{\partial J}{\partial \text{net}_k} = -\frac{\partial J}{\partial z_k} \cdot \frac{\partial z_k}{\partial \text{net}_k} = (t_k - z_k) f'(\text{net}_k)$$

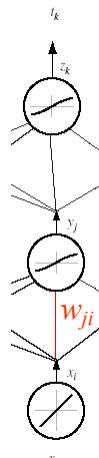
$$\frac{\partial \text{net}_k}{\partial w_{kj}} = y_j$$

$$\boxed{\Delta w_{kj} = -\eta \frac{\partial J}{\partial w_{kj}} = \eta \delta_k y_j = \eta (t_k - z_k) f'(\text{net}_k) y_j}$$

Network Learning

Backpropagation (2)

- Input-to-hidden weights w_{ji}



$$\frac{\partial J}{\partial w_{ji}} = \frac{\partial J}{\partial y_j} \cdot \frac{\partial y_j}{\partial \text{net}_j} \cdot \frac{\partial \text{net}_j}{\partial w_{ji}} = \frac{\partial J}{\partial y_j} \cdot f'(\text{net}_j) x_i$$

$$\begin{aligned} \frac{\partial J}{\partial y_j} &= \frac{\partial}{\partial y_j} \left[\frac{1}{2} \sum_{k=1}^c (t_k - z_k)^2 \right] = -\sum_{k=1}^c (t_k - z_k) \frac{\partial z_k}{\partial y_j} \\ &= -\sum_{k=1}^c (t_k - z_k) \frac{\partial z_k}{\partial \text{net}_k} \frac{\partial \text{net}_k}{\partial y_j} = -\sum_{k=1}^c (t_k - z_k) \underbrace{f'(\text{net}_k) w_{kj}}_{\delta_k} \end{aligned}$$

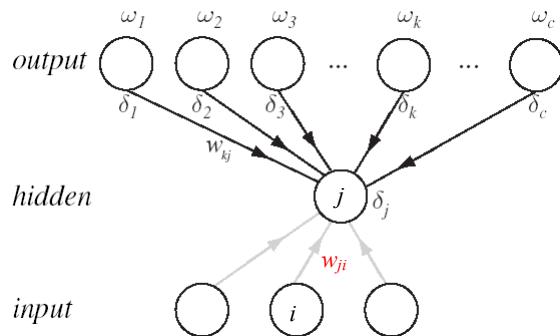
$$\text{define } \delta_j = \frac{\partial J}{\partial y_j} \cdot f'(\text{net}_j) = f'(\text{net}_j) \sum_{k=1}^c w_{kj} \delta_k$$

Network Learning

Backpropagation (3)

- Input-to-hidden weights w_{ji}

$$\Delta w_{ji} = -\eta \frac{\partial J}{\partial w_{ji}} = \eta \delta_j x_i = \eta \left[\sum_{k=1}^c w_{kj} \delta_k \right] f'(net_j) x_i$$



Network Learning

Backpropagation (4)

- Initialization
 - Small random non-zero values... why?
- Stochastic version

```

Initialize {nH, w, θ, η, k ← 0}
do {
  k ← k+1
  xk ← randomly chosen pattern
  wkj ← wkj + ηδkyj
  wji ← wji + ηδjxi
} until ( ||∇J(w)|| < θ)
  
```

- Batch version: $J = \sum_{p=1}^n J_p$

Network Learning

Backpropagation (5)

- Stopping criterion
 - stop training at a minimum of the error on the test or validation set

