

# **Part 3: Deadlocks**

# Overview

---

- ❖ Resources
- ❖ Why do deadlocks occur?
- ❖ Dealing with deadlocks
  - Ignoring them: ostrich algorithm
  - Detecting & recovering from deadlock
  - Avoiding deadlock
  - Preventing deadlock

# Resources

---

- ❖ Resource: something a process uses
  - Usually limited (at least somewhat)
- ❖ Examples of computer resources
  - Printers
  - Semaphores / locks
  - Tables (in a database)
- ❖ Processes need access to resources in reasonable order
- ❖ Two types of resources:
  - Preemptable resources: can be taken away from a process with no ill effects
  - Nonpreemptable resources: will cause the process to fail if taken away


# When do deadlocks happen?

- ❖ Suppose

- Process 1 holds resource A and requests resource B
- Process 2 holds B and requests A
- Both can be blocked, with neither able to proceed

- ❖ Deadlocks occur when ...

- Processes are granted exclusive access to devices or software constructs (resources)
- Each deadlocked process needs a resource held by another deadlocked process



# Using resources

- ❖ Sequence of events required to use a resource
  - Request the resource
  - Use the resource
  - Release the resource
- ❖ Can't use the resource if request is denied
  - Requesting process has options
    - Block and wait for resource
    - Continue (if possible) without it: may be able to use an alternate resource
    - Process fails with error code
  - Some of these may be able to prevent deadlock...

# What is a deadlock?

---


- ❖ Formal definition:  
**“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”**
- ❖ Usually, the event is release of a currently held resource
- ❖ In deadlock, none of the processes can
  - Run
  - Release resources
  - Be awakened

# Four conditions for deadlock

---

- ◆ Mutual exclusion
  - Each resource is assigned to at most one process
- ◆ Hold and wait
  - A process holding resources can request more resources
- ◆ No preemption
  - Previously granted resources cannot be forcibly taken away
- ◆ Circular wait
  - There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain

# Resource allocation graphs



- ❖ Resource allocation modeled by directed graphs
- ❖ Example 1:
  - Resource R assigned to process A
- ❖ Example 2:
  - Process B is requesting / waiting for resource S
- ❖ Example 3:
  - Process C holds T, waiting for U
  - Process D holds U, waiting for T
  - C and D are in deadlock!


# Dealing with deadlock

- ✿ How can the OS deal with deadlock?
  - Ignore the problem altogether!
    - Hopefully, it'll never happen...
  - Detect deadlock & recover from it
  - Dynamically avoid deadlock
    - Careful resource allocation
  - Prevent deadlock
    - Remove at least one of the four necessary conditions
- ✿ We'll explore these tradeoffs

# Getting into deadlock


A

Acquire R  
Acquire S  
Release R  
Release S







B

Acquire S  
Acquire T  
Release S  
Release T



C

Acquire T  
Acquire R  
Release T  
Release R



# Not getting into deadlock...

- ♦ Many situations may result in deadlock (but don't have to)
  - In previous example, A could release R before C requests R, resulting in no deadlock
  - Can we always get out of it this way?
- ♦ Find ways to:
  - Detect deadlock and reverse it
  - Stop it from happening in the first place

# The Ostrich Algorithm


---

- ❖ Pretend there's no problem
- ❖ Reasonable if
  - Deadlocks occur very rarely
  - Cost of prevention is high
- ❖ UNIX and Windows take this approach
  - Resources (memory, CPU, disk space) are plentiful
  - Deadlocks over such resources rarely occur
  - Deadlocks typically handled by rebooting
- ❖ Trade off between convenience and correctness

# Detecting deadlocks using graphs

- Process holdings and requests in the table and in the graph (they're equivalent)
- Graph contains a cycle  $\Rightarrow$  deadlock!
  - Easy to pick out by looking at it (in this case)
  - Need to mechanically detect deadlock
- Not all processes are deadlocked (A, C, F not in deadlock)

Process	Holds	Wants
A	R	S
B		T
C		S
D	U	S,T
E	T	V
F	W	S
G	V	U



# Deadlock detection algorithm

- ♦ General idea: try to find cycles in the resource allocation graph
- ♦ Algorithm: depth-first search at each node
  - Mark arcs as they're traversed
  - Build list of visited nodes
  - If node to be added is already on the list, a cycle exists!
- ♦ Cycle  $\Rightarrow$  deadlock

```
For each node N in the graph {  
    Set L = empty list  
    unmark all arcs  
    Traverse (N,L)  
}  
If no deadlock reported by now,  
there isn't any  
  
define Traverse (C,L) {  
    If C in L, report deadlock!  
    Add C to L  
    For each unmarked arc from C {  
        Mark the arc  
        Set A = arc destination  
        /* NOTE: L is a  
           local variable */  
        Traverse (A,L)  
    }  
}
```

# Resources with multiple instances

- ❖ Previous algorithm only works if there's one instance of each resource
- ❖ If there are multiple instances of each resource, we need a different method
  - Track current usage and requests for each process
  - To detect deadlock, try to find a scenario where all processes can finish
  - If no such scenario exists, we have deadlock

# Deadlock detection algorithm

	A	B	C	D
Avail	2	3	0	1
Process	A	B	C	D
1	0	3	0	0
2	1	0	1	1
3	0	2	1	0
4	2	2	3	0

Hold

Process	A	B	C	D
1	3	2	1	0
2	2	2	0	0
3	3	5	3	1
4	0	4	1	1

Want

```
current=avail;
for (j = 0; j < N; j++) {
    for (k=0; k<N; k++) {
        if (finished[k])
            continue;
        if (want[k] < current) {
            finished[k] = 1;
            current += hold[k];
            break;
        }
        if (k==N) {
            printf "Deadlock!\n";
            // finished[k]==0 means
            process
            // is in the deadlock
            break;
        }
    }
}
```

Note: want[j],hold[j],current,avail are arrays!

# Recovering from deadlock

- ❖ Recovery through **preemption**

- Take a resource from some other process
  - Depends on nature of the resource and the process


- ❖ Recovery through **rollback**

- Checkpoint a process periodically
  - Use saved state to restart the process if it's in deadlock
  - May present a problem if the process affects lots of “external” things

- ❖ Recovery through **killing processes**

- Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle
  - Other processes can get its resources
  - Try to choose a process that can be rerun from the start
    - Pick one that hasn't run too far already

# Resource trajectories



Two process resource trajectories

# Safe and unsafe states

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3

	Has	Max
A	3	9
B	4	4
C	2	7

Free: 1

	Has	Max
A	3	9
B	0	-
C	2	7

Free: 5

	Has	Max
A	3	9
B	0	-
C	7	7

Free: 0

	Has	Max
A	3	9
B	0	-
C	0	-

Free: 7

Demonstration that the first state is safe

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3

	Has	Max
A	4	9
B	2	4
C	2	7

Free: 2

	Has	Max
A	4	9
B	4	4
C	2	7

Free: 0

	Has	Max
A	4	9
B	0	-
C	2	7

Free: 4

Demonstration that the second state is unsafe

# Banker's Algorithm for a single resource

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

Any sequence finishes

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

C,B,A,D finishes

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

Deadlock (unsafe state)

- ✿ Bankers' algorithm: before granting a request, ensure that a **sequence** exists that will allow all processes to complete
  - Use previous methods to find such a sequence
  - If a sequence exists, allow the requests
  - If there's no such sequence, deny the request
- ✿ Can be slow: must be done on each request!

# Banker's Algorithm for multiple resources

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

$$\begin{aligned}E &= (6342) \\P &= (5322) \\A &= (1020)\end{aligned}$$

Example of banker's algorithm with multiple resources

# Preventing deadlock

---

- ❖ Deadlock can be completely prevented!
- ❖ Ensure that at least one of the conditions for deadlock never occurs
  - Mutual exclusion
  - Circular wait
  - Hold & wait
  - No preemption
- ❖ Not always possible...

# Eliminating mutual exclusion


- ❖ Some devices (such as printer) can be **spooled**
  - Only the printer daemon uses printer resource
  - This eliminates deadlock for printer
- ❖ Not all devices can be spooled
- ❖ Principle:
  - Avoid assigning resource when not absolutely necessary
  - As few processes as possible actually claim the resource

# Attacking “hold and wait”

- ❖ Require processes to **request** resources **before** starting
  - A process never has to wait for what it needs
- ❖ This can present problems
  - A process may not know required resources at start of run
  - This also ties up resources other processes could be using
    - Processes will tend to be conservative and request resources they might need
- ❖ Variation: a process must **give up** all resources before making a new request
  - Process is then granted all prior resources as well as the new ones
  - Problem: what if someone grabs the resources in the meantime—how can the process save its state?


# Attacking “no preemption”

- ❖ This is not usually a viable option
- ❖ Consider a process given the printer
  - Halfway through its job, take away the printer
    - Confusion ensues!
- ❖ May work for some resources
  - Forcibly take away memory pages, suspending the process
  - Process may be able to resume with no ill effects



# Attacking “circular wait”

- ◆ Assign an order to resources
- ◆ Always acquire resources in numerical order
  - Need not acquire them all at once!
- ◆ Circular wait is prevented
  - A process holding resource n can't wait for resource m if  $m < n$
  - No way to complete a cycle
    - Place processes above the highest resource they hold and below any they're requesting
    - All arrows point up!



# Deadlock prevention: summary

---

- ❖ Mutual exclusion
  - Spool everything
- ❖ Hold and wait
  - Request all resources initially
- ❖ No preemption
  - Take resources away
- ❖ Circular wait
  - Order resources numerically

# Example: two-phase locking

- ❖ Phase One
  - Process tries to lock all data it needs, one at a time
  - If needed data found locked, start over
  - (no real work done in phase one)
- ❖ Phase Two
  - Perform updates
  - Release locks
- ❖ Note similarity to requesting all resources at once
- ❖ This is often used in databases
- ❖ It avoids deadlock by eliminating the “hold-and-wait” deadlock condition

# “Non-resource” deadlocks

---

- ❖ Possible for two processes to deadlock
  - Each is waiting for the other to do some task
- ❖ Can happen with semaphores
  - Each process required to do a `down()` on two semaphores (`mutex` and another)
  - If done in wrong order, deadlock results
- ❖ Semaphores could be thought of as resources...

# Starvation

---

- ❖ Algorithm to allocate a resource
  - Give the resource to the **shortest** job first
- ❖ Works great for multiple short jobs in a system
- ❖ May cause long jobs to be postponed indefinitely
  - Even though not blocked
- ❖ Solution
  - First-come, first-serve policy
- ❖ Starvation can lead to deadlock
  - Process starved for resources can be holding resources
  - If those resources aren't used and released in a timely fashion, **shortage** could lead to deadlock

# Livelock

---

- ❖ Sometimes, processes can still run, but not make progress
- ❖ Example: two processes want to use resources A and B
  - P0 gets A, P1 gets B
  - Each realizes that a deadlock will occur if they proceed as planned!
  - P0 drops A, P1 drops B
  - P0 gets B, P1 gets A
  - Same problem as before
  - This can go on for a very long time...
- ❖ Real-world example: Ethernet transmission collisions
  - If there's a “collision” on the wire, wait and try again
  - Multiple processes waited the exact same amount of time...