

COEN281 -- Introduction to Pattern Recognition and Data Mining

Lecture 8: Neural Networks (Part 2)

Instructor: Dr. Giovanni Seni
GSeni@scu.edu

***Department of Computer Engineering
Santa Clara University***

Fall/18

Syllabus

Week 1	Introduction; R (Ch.1)
Week 2	Bayesian Decision Theory (Ch.2; DHS: 2.1-2.6, 2.9) Parameter Estimation (DHS: 3.1-3.4)
Week 3	Linear Discriminant Functions (Ch.3&4; DHS: 3.8.2, 5.1-5.8) Regularization (Ch.6; SE: Ch.3)
Week 4	Neural Networks (DHS: 6.1-6.6, 6.8); Deep Learning
Week 5	Support Vector Machines (Ch.9)
Week 6	Decision Trees (Ch. 8.1; DHS: 8.3; Ch 2 SE)
Week 7	Ensemble Methods (Ch. 8.2; SE: Ch 4, 5)
Week 8	Clustering (Ch. 10; DHS: 10.6, 10.7) Clustering (DHS: 10.9); How many clusters are there? (DHS: 10.10)
Week 9	Non-metric: Association Rules Collaborative Filtering
Week 10	Text Retrieval; Other topics

Overview

- Error Surfaces
 - Minimums and Plateaus
- Practical Considerations
 - Momentum
 - Scaling input
 - Target values
 - Number of hidden units
 - Learning rate
 - Minibatch

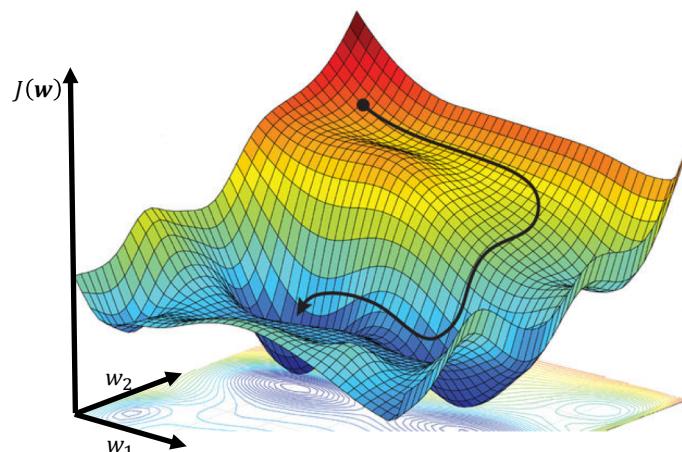
COEN281

GSeni@scu.edu

3

Error Surfaces

Minimas



- Single global min exists; **nice but not essential!**

COEN281

GSeni@scu.edu

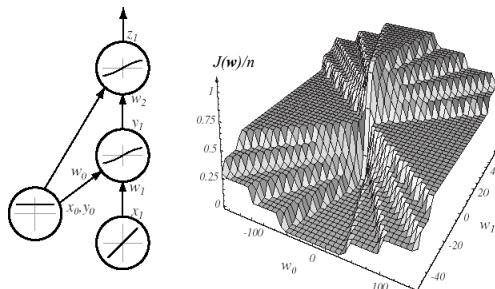
4

Error Surfaces

Plateaus

Plateaus – regions where weight change does not lead to a change in error

- Training can be slow
- Example: separable 2-class problem:
 - At $w \approx 0$ error is high and slope is large



COEN281

GSeni@scu.edu

5

Practical Considerations

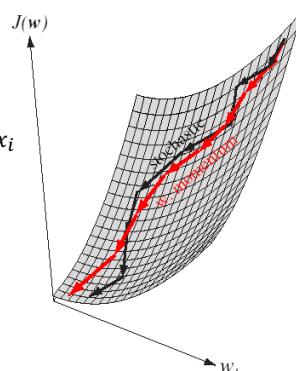
Momentum

- At plateaus, $J'(\mathbf{w}) \approx 0$ and learning can be very slow
- Physics' notion: keep moving unless acted upon by outside forces
- Modified learning rule ($\alpha \approx 0.9$):

$$w_{ji} = w_{ji} + \Delta w \quad \text{where } \Delta w = -\eta \nabla J(\mathbf{w}) = \eta \delta_j x_i$$

↓

$$\Delta w_t = \alpha \cdot \Delta w_{t-1} + (1 - \alpha) \cdot \eta \nabla J(\mathbf{w})$$



COEN281

GSeni@scu.edu

6

Practical Considerations

Scaling Input

- If one attribute is much larger than another attribute, the weights will be adjusted to represent such differences; that is not desirable
 - Example: mass (grams) and length (meters) of a fish
- Solution:
 - Standardize all features before training
 - E.g., “whitening” transform: $x_{new} = (x - \mu) / \sigma$
 - The mean of each feature is set to zero and the variance to 1.0
 - Also, convergence is usually faster if the average of each input variable over the training set is close to zero

COEN281

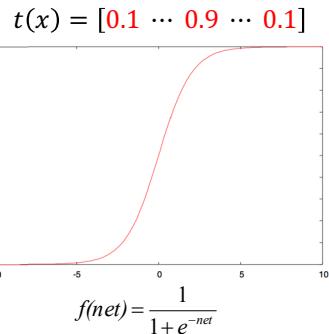
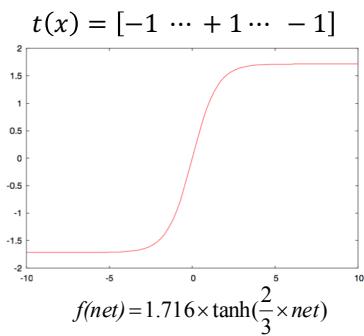
GSeni@scu.edu

7

Practical Considerations

Target Values

- One-of- c representation



- Saturation values should be avoided

COEN281

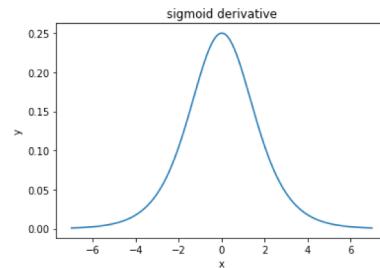
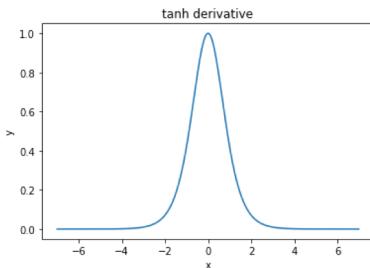
GSeni@scu.edu

8

Practical Considerations

Tanh vs Logistic

- \tanh has larger derivatives

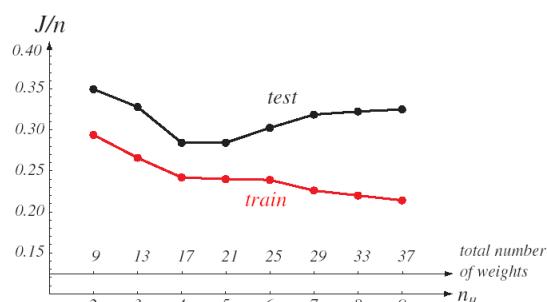


- 0-centered output is also desirable (see scaling)

Practical Considerations

Number of Hidden Units

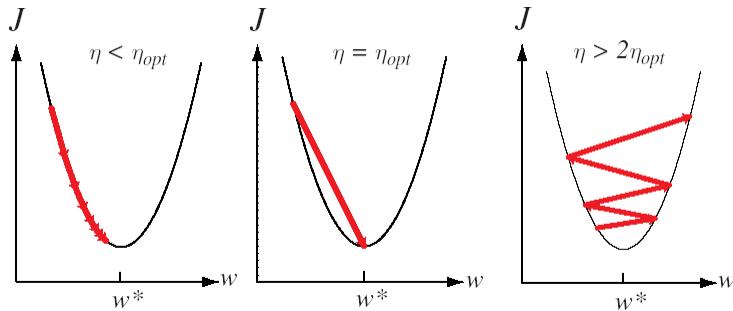
- n_H governs the expressive power of the net
 - Number of weights \approx number of degrees of freedom
 - Should never have more weights than training patterns!
 - Rule of thumb: $n/10$



Practical Considerations

Learning Rate

- In principle, if small enough (e.g., $\eta \approx 0.001$), it only determines speed of convergence



- A principled method exists for rapid and uniform learning

$$\eta(w_k) = 1/J''(w_k)$$

COEN281

GSeni@scu.edu

11

Practical Considerations

Optimal Learning Rate

- 1-Dimensional Quadratic Example: $J(w) = aw^2 + bw + c$

$$\begin{aligned} \Rightarrow J'(w) &= 2aw + b & J'(w) = 0 \Rightarrow 2aw + b = 0 \\ J''(w) &= 2a & \hat{w} = -b/2a \end{aligned}$$

- Now imagine that we are starting at some arbitrary initial point w_0 and we want to get to \hat{w}

- Step $\Delta w = -\eta J'(w)$
 - Best η : get to \hat{w} in one step – i.e., $\Delta w = \hat{w} - w_0$

$$\Rightarrow -\frac{b}{2a} - w_0 = \frac{2aw_0 + b}{2a} = \eta \cdot (2aw_0 + b)$$

$$\Rightarrow \eta = \frac{1}{2a} = \frac{1}{J''(w_0)}$$

COEN281

GSeni@scu.edu

12

Practical Considerations

Learning Rate Decay

- Constant

(requires schedule for piecewise constant, tricky)

- Useful hack

Keep learning rate constant until no improvement on validation data, then reduce rate.

- Polynomial decay

Canonical choice for convex solvers $\eta(t) = \frac{\alpha}{\sqrt{\beta+t}}$

- Exponential decay

not recommended

Practical Considerations

Learning Rate Decay (2)

- (locally) Adaptive

$$\eta_{ij}(t) = \frac{\eta_0}{\sqrt{K + \sum_{t'=t-\tau}^t J''_{ij}(t')}}$$

- For directions with large gradient, decrease learning rate aggressively to avoid instability
- If gradients start vanishing, learning rate decrease reduces, too

Practical Considerations

Minibatch

- The frequent updates of Stochastic Gradient Descent (SGD) can result in a noisy gradient signal
 - Weights can jump around ==> hard to settle on an error minimum

Aggregate gradients over a few instances before updating

- Reduces variance in gradients
- Better for vectorization (GPUs)
- Large minibatch may need lots of memory (and may slow updates)
- GPU considerations
 - 256GPUs and 20 samples/GPU ==> mini batch of 5120