

## COEN281 -- Introduction to Pattern Recognition and Data Mining

### Lecture 10: Support Vector Machines

Instructor: Dr. Giovanni Seni  
[GSeni@scu.edu](mailto:GSeni@scu.edu)

*Department of Computer Engineering  
Santa Clara University*

Fall/18

## Syllabus

|         |   |
|---------|---|
| Week 1  | Introduction; R (Ch.1)  |
| Week 2  | Bayesian Decision Theory (Ch.2; DHS: 2.1-2.6, 2.9)<br>Parameter Estimation (DHS: 3.1-3.4)                 |
| Week 3  | Linear Discriminant Functions (Ch.3&4; DHS: 3.8.2, 5.1-5.8)<br>Regularization (Ch.6; SE: Ch.3)            |
| Week 4  | Neural Networks (DHS: 6.1-6.6, 6.8); Deep Learning  |
| Week 5  | <b>Support Vector Machines</b> (Ch.9)   |
| Week 6  | Decision Trees (Ch. 8.1; DHS: 8.3; Ch 2 SE)   |
| Week 7  | Ensemble Methods (Ch. 8.2; SE: Ch 4, 5)   |
| Week 8  | Clustering (Ch. 10; DHS: 10.6, 10.7)<br>Clustering (DHS: 10.9); How many clusters are there? (DHS: 10.10) |
| Week 9  | Non-metric: Association Rules<br>Collaborative Filtering  |
| Week 10 | Text Retrieval; Other topics  |

## Overview

- Introduction
  - Geometry of hyperplanes revisited
  - Generalized linear discriminant functions revisited
  - Support Vectors
- SVM Training
  - Overview
  - Optimization approach
  - Non-separable (overlap) case
- The Kernel trick
- SVM as a “penalized” method
- SVM for regression

## Introduction

### Approaches to Building Classifiers

- 1) *Class-conditional* (“generative”) approach
  - $p_{\theta_j}(x|\omega_j)$  are modeled explicitly;  $\hat{\theta}_j$  are estimated via ML. Combined with estimates of  $P(\omega_j)$  are inverted via Bayes rule to arrive at  $P(\omega_j|x)$
  - E.g., **LDA**
- 2) *Regression* approach
  - $P(\omega_j|x)$  are modeled explicitly
  - E.g., **Logistic regression, NNs**
- 3) *Discriminative* approach
  - Try to model the decision boundary directly
  - E.g., **SVMs**

## Introduction

### Geometry of Hyperplanes Revisited

- Equation  $g(\mathbf{y})=0$  defines surface that separates points assigned to the category  $\omega_1$  from points assigned to the category  $\omega_2$

–  $g(\mathbf{y})$  linear  $\Rightarrow$  surface is a *hyperplane*  $H$


– Consider  $\mathbf{y}_1$  and  $\mathbf{y}_2$  both on the decision surface:

$$\mathbf{a}'\mathbf{y}_1 + a_0 = \mathbf{a}'\mathbf{y}_2 + a_0$$

$$\text{or } \mathbf{a}'(\mathbf{y}_1 - \mathbf{y}_2) = 0$$

$\Rightarrow \mathbf{a}$  is normal to any vector lying in the hyperplane

– Orientation of  $H$  is determined by  $\mathbf{a}$



## Introduction

### Geometry of Hyperplanes Revisited (2)

- $g(\mathbf{y}) \propto$  distance from  $\mathbf{y}$  to  $H$


– Express  $\mathbf{y}$  as  $\mathbf{y} = \mathbf{y}_p + b \frac{\mathbf{a}}{\|\mathbf{a}\|}$

because  $g(\mathbf{y}_p) = 0$

$$g(\mathbf{y}) = \mathbf{a}'\mathbf{y} + a_0 = g(\mathbf{y}_p) + b \frac{\mathbf{a}'\mathbf{a}}{\|\mathbf{a}\|}$$

$$= b \|\mathbf{a}\|$$

$$\Rightarrow b = \frac{g(\mathbf{y})}{\|\mathbf{a}\|}$$




- Also,  $d(0, H) = a_0 / \|\mathbf{a}\|$  -i.e., location of  $H$  is determined by  $a_0$

## Introduction

### Generalized Linear Discriminant


- XOR



## Introduction

### Which Hyperplane to Use?


- Infinite number of separating hyperplanes... How to pick?



## Maximal Margin Hyperplane

### Support Vectors

- Separating hyperplane that is furthest from the training observations



## Maximal Margin Hyperplane

### Support Vectors (2)

- SVs: (transformed) training patterns that define the optimal separating hyperplane
  - Equidistant from the hyperplane
  - Informally speaking, they are the most "informative" for classification
- Points well inside their class boundary play a very small role in shaping the boundary
  - This is an attractive property
- In contrast, the decision boundary in LDA is determined by  $\mu_i, \mu_j, \Sigma$ 
  - In this regard, logistic regression is more similar to the support vector classifier

## SVM Training

### Overview

- Assume each pattern  $\mathbf{x}_k$  has been transformed to  $\mathbf{y}_k = \varphi(\mathbf{x}_k)$
- Class encoding: let  $z_k = \pm 1$ , according to whether  $\mathbf{x}_k$  is in  $\omega_1$  or  $\omega_2$
- A separating hyperplane ensures  $z_k g(\mathbf{y}_k) \geq 1$
- Since  $d(\mathbf{y}, H) = |g(\mathbf{y})| / \|\mathbf{a}\|$ , and assuming a positive margin  $b$  exists:
$$\frac{z_k g(\mathbf{y}_k)}{\|\mathbf{a}\|} \geq b \quad \forall k$$
- Goal: find weight vector  $\mathbf{a}$  that maximizes  $b$ 
  - The larger the margin, the better the generalization of the classifier

## SVM Training

### Overview (2)

- Ensure solution uniqueness by imposing constraint  $b \cdot \|\mathbf{a}\| = 1$
- The optimization problem

$$\begin{aligned} & \max_{\mathbf{a}, a_0, \|\mathbf{a}\|=1} b \\ & \text{subject to } z_k g(\mathbf{y}_k) \geq b, \quad k = 1, \dots, N \end{aligned}$$

- Or more conveniently rephrased as

$$\begin{aligned} & \min_{\mathbf{a}, a_0} \|\mathbf{a}\| \\ & \text{subject to } z_k g(\mathbf{y}_k) \geq 1, \quad k = 1, \dots, N \end{aligned}$$

- Convex optimization problem: quadratic criterion, linear inequality constraints

## SVM Training


### Method of Lagrange Multipliers

- Relax the constraint with parameter Lambda, and put it into the objective function

$$L(\mathbf{a}, \lambda) = \frac{1}{2} \|\mathbf{a}\|^2 - \sum_{k=1}^N \lambda_k [z_k g(\mathbf{y}_k) - 1]$$

- the  $\lambda_k$  's can be considered as a penalty cost of violating the constraint  $z_k g(\mathbf{y}_k) \geq 1$

- We seek to maximize  $L$  with respect to  $\lambda$ , and minimize with respect to  $\mathbf{a}$



## SVM Training

### Method of Lagrange Multipliers (2)

- At the saddle point,  $\partial L / \partial a_0 = 0$  and  $\partial L / \partial \mathbf{a} = 0$

$$-\partial L / \partial a_0 = 0 \Rightarrow \sum_{k=1}^N \lambda_k^* \cdot z_k = 0$$

$$-\partial L / \partial \mathbf{a} = 0 \Rightarrow \partial L / \partial \mathbf{a} = \mathbf{a} - \sum_{k=1}^N \lambda_k^* \cdot z_k \cdot \mathbf{y}_k = 0 \Rightarrow \mathbf{a}^* = \sum_{k=1}^N \lambda_k^* \cdot z_k \cdot \mathbf{y}_k$$


$\Rightarrow$  The optimal hyperplane is a linear combination of the (augmented) training patterns

- Thus,  $\hat{g}(\mathbf{y}) = \hat{\mathbf{a}}^t \mathbf{y} + \hat{a}_0 = \hat{a}_0 + \sum_{k=1}^N \lambda_k^* \cdot z_k \cdot \mathbf{y} \cdot \mathbf{y}_k$

- Can be shown  $\lambda_k$  is non-zero if and only if  $\mathbf{y}_k$  is a support vector

## SVM Training

### Non-separable Case & Noisy Data



COEN281

GSeni@scu.edu

15

## SVM Training

### Soft-margin Classifier

- Allow for some points to be on the wrong side of the margin


- Define “slack” variables

$$\xi = (\xi_1, \xi_2, \dots, \xi_N)$$

- Rephrased problem:

$$\min_{\mathbf{a}, a_0} \|\mathbf{a}\|$$

$$\text{subject to } \begin{cases} z_k g(\mathbf{y}_k) \geq 1 - \xi_k & \forall k, \\ \xi_k \geq 0, \quad \sum_{k=1}^N \xi_k \leq C \end{cases}$$



- $C$  is tuned via cross-validation (bias-variance trade-off)

COEN281

GSeni@scu.edu

16

## The Kernel Trick

- $\mathbf{y} = \varphi(\mathbf{x})$  is involved only through the inner products  $\mathbf{y} \cdot \mathbf{y}_k$
- Given a suitable “kernel” function  $K(\mathbf{x}, \mathbf{x}_k) = \mathbf{y} \cdot \mathbf{y}_k$ , we don’t need  $\varphi(\mathbf{x})$  at all
  - $K$  is a symmetric positive (semi-) definite function
- Popular choices of  $K$ :

**$d$ -th Degree polynomials:**  $K(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x} \cdot \mathbf{x}')^d$

**Radial basis:**  $K(\mathbf{x}, \mathbf{x}') = \exp\left(-\|\mathbf{x} - \mathbf{x}'\|^2/c\right)$

**Neural network:**  $K(\mathbf{x}, \mathbf{x}') = \tanh(\kappa_1 \mathbf{x} \cdot \mathbf{x}' + \kappa_2)$

## The Kernel Trick (2)

- Example: 2-th degree polynomials  $K(\mathbf{x}, \mathbf{x}') = (1 + \mathbf{x} \cdot \mathbf{x}')^2$ ,  $\mathbf{x} \in \mathbb{R}^2$

$$\begin{aligned} \rightarrow K(\mathbf{x}, \mathbf{x}') &= (1 + \mathbf{x} \cdot \mathbf{x}')^2 \\ &= (1 + x_1 x'_1 + x_2 x'_2)^2 \\ &= 1 + 2x_1 x'_1 + 2x_2 x'_2 + (x_1 x'_1)^2 + (x_2 x'_2)^2 + 2x_1 x'_1 x_2 x'_2 \end{aligned}$$

If we choose  $\mathbf{y} = \varphi(\mathbf{x}) = \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ x_1^2 \\ x_2^2 \\ \sqrt{2}x_1 x_2 \end{bmatrix}$ , then  $K(\mathbf{x}, \mathbf{x}') = \mathbf{y} \cdot \mathbf{y}'$

## The Kernel Trick (3)

- So we write our discriminant function as


$$\hat{g}(\mathbf{x}) = \hat{a}_0 + \sum_{k=1}^N \hat{\lambda}_k \cdot z_k \cdot \mathbf{y} \cdot \mathbf{y}_k = \hat{a}_0 + \sum_{i \in SV} \alpha_i \cdot K(\mathbf{x}, \mathbf{x}_i)$$

- Classifier's complexity is characterized by the number of support vectors, not the dimensionality of the transformed space
  - Less prone to problems of overfitting? No! (See ESL 12.3.4)
  - Looking at  $K(\mathbf{x}, \mathbf{x})$ :
    - Kernel cannot adapt itself to concentrate on subspaces – i.e., lacks ability to ignore redundant terms
    - A higher than necessary degree polynomial will do worse

## SVM as a Penalized Method

- Consider the (“hinge”) loss
$$L(z, g) = [1 - zg]_+$$
- Hinge and log-likelihood have similar tails
  - Zero penalty to points well inside their margin
- With  $g(\mathbf{x}) = \mathbf{a}'\varphi(\mathbf{x}) + a_0$ , the optimization problem

$$\min_{\mathbf{a}, a_0} \sum_{k=1}^N L(z_k, g(\mathbf{x}_k)) + \gamma \|\mathbf{a}\|^2$$




can be shown to have the same solution as Slide 14 ( $\gamma = 1/C$ )

## SVM for Regression

- Consider the (“ $\varepsilon$ -insensitive”) loss

$$L_\varepsilon(z, g) = \begin{cases} 0 & \text{if } |z - g| < \varepsilon \\ |z - g| - \varepsilon & \text{otherwise} \end{cases}$$

- Points on the correct side of the boundary  $\approx$  points with small residual  
 $\Rightarrow$  ignored in the optimization



- Solution to  $\min_{\mathbf{a}, a_0} \sum_{k=1}^N L_\varepsilon(z_k, g(\mathbf{x}_k)) + \frac{\gamma}{2} \|\mathbf{a}\|^2$  has the form

$$\hat{\mathbf{a}} = \sum_{k=1}^N (\hat{\lambda}_k^* - \hat{\lambda}_k) \mathbf{x}_k \quad \hat{g}(\mathbf{x}) = a_0 + \sum_{k=1}^N (\hat{\lambda}_k^* - \hat{\lambda}_k) \mathbf{x} \cdot \mathbf{x}_k$$

## SVM for Regression (2)

- Where  $\hat{\lambda}_k, \hat{\lambda}_k^*$  solve the quadratic programming problem

$$\min_{\lambda_k, \lambda_k^*} \varepsilon \sum_{k=1}^N (\lambda_k^* - \lambda_k) - z_k \sum_{k=1}^N (\lambda_k^* - \lambda_k) + \frac{1}{2} \sum_{j,k=1}^N (\lambda_k^* - \lambda_k)(\lambda_j^* - \lambda_j) \mathbf{x}_k \cdot \mathbf{x}_j$$

Subject to the constraints

$$0 \leq \lambda_k, \lambda_k^* \leq 1/\gamma \quad \sum_{k=1}^N (\lambda_k^* - \lambda_k) = 0 \quad \lambda_k \lambda_k^* = 0$$

- Nature of constraints imply typically only a subset of  $(\lambda_k^* - \lambda_k)$  are non-zero
- Criterion meta-parameters:  $\varepsilon$  and  $\gamma$