

## **COEN281 -- Introduction to Pattern Recognition and Data Mining**

### **Lecture 11: Classification & Regression Trees**

Instructor: Dr. Giovanni Seni

*GSeni@scu.edu*

***Department of Computer Engineering  
Santa Clara University***

Fall/18

## **Syllabus**

Week 1	Introduction; R (Ch.1)	
Week 2	Bayesian Decision Theory (Ch.2; DHS: 2.1-2.6, 2.9) Parameter Estimation (DHS: 3.1-3.4)	
Week 3	Linear Discriminant Functions (Ch.3&4; DHS: 3.8.2, 5.1-5.8) Regularization (Ch.6; SE: Ch.3)	
Week 4	Neural Networks (DHS: 6.1-6.6, 6.8); Deep Learning	Predictive Learning
Week 5	Support Vector Machines (Ch.9)	
Week 6	<b>Decision Trees</b> (Ch. 8.1; DHS: 8.3; Ch 2 SE)	
Week 7	Ensemble Methods (Ch. 8.2; SE: Ch 4, 5)	
Week 8	Clustering (Ch. 10; DHS: 10.6, 10.7) Clustering (DHS: 10.9); How many clusters are there? (DHS: 10.10)	
Week 9	Non-metric: Association Rules Collaborative Filtering	
Week 10	Text Retrieval; Other topics	

## Overview

---

- Motivation
- Predictive Learning Example
- Tree Induction
  - Growing
  - Pruning
- Key Features
- *Surrogate* splits
- Limitations

## Motivation

---

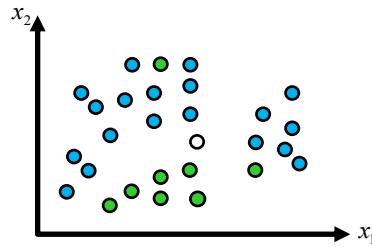
What's "tricky" with LDA, LogReg, NNets, SVMs?

- Handling of categorical features
- Sensitivity to variable scaling
- Sensitivity to bad  $x_j$  - distributions (e.g., outliers)
- Handling of missing values
- Interpretability (for NNets and SVMs)

## Predictive Learning Example

- A simple data set

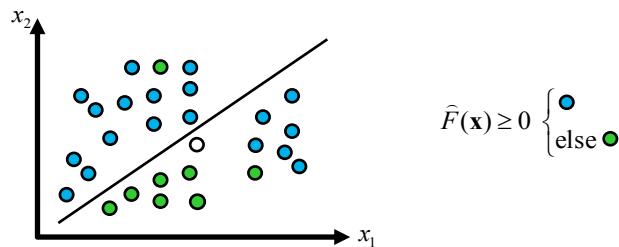
Attribute-1 ( $x_1$ )	Attribute-2 ( $x_2$ )	Class
1.0	2.0	blue
2.0	1.0	green
...	...	...
4.5	3.5	?



- What is the class of new point  $\circ$ ?
- Many approaches... no method is universally better; try several / use committee

## Predictive Learning Example (2)

- Ordinary Linear Regression (OLR)

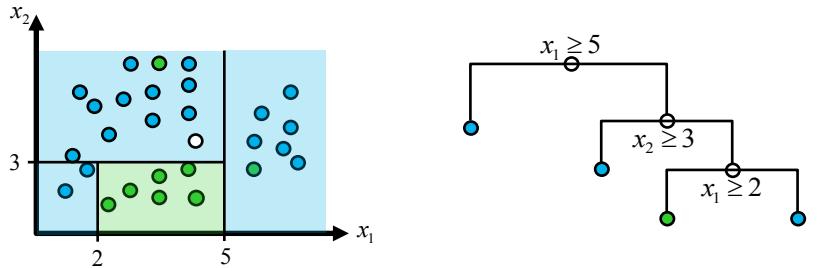


- Model:  $\hat{F}(\mathbf{x}) = a_0 + \sum_{j=1}^n a_j x_j \quad \forall (a_0, \mathbf{a})$   $\Rightarrow$  Not flexible enough
- Accuracy (on training data): 74.07%

## Decision Trees

### Classification Example

- Discrete y

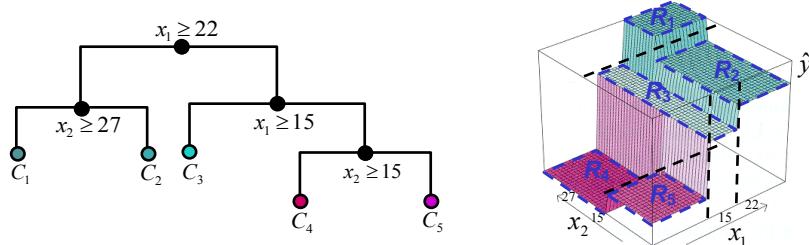


- Model:  $\hat{F}(\mathbf{x}) = \sum_{m=1}^M \hat{c}_m I_{\hat{R}_m}(\mathbf{x})$        $\{\hat{R}_m\}_{m=1}^M$  = Subregions of input variable space
- Accuracy (on training data): 92.59%

## Decision Trees

### Regression Example

- Continuous y



- Model:  $\hat{y} = T(\mathbf{x}) = \sum_{m=1}^M \hat{c}_m I_{\hat{R}_m}(\mathbf{x})$   
where  $I_A(\mathbf{x}) = 1$  if  $\mathbf{x} \in A$ , 0 otherwise

## Tree Induction

### Theory Overview

- Model:  $\hat{F}(\mathbf{x}) = \sum_{m=1}^M \hat{c}_m I_{\hat{R}_m}(\mathbf{x})$
- Score criterion:
  - Regression – least squares – i.e.,  $L(y, \hat{y}) = (y - \hat{y})^2$

$$\{\hat{c}_m, \hat{R}_m\}_1^M = \arg \min_{\{c_m, R_m\}_1^M} \sum_{i=1}^N \left[ y_i - \sum_{m=1}^M c_m I_{R_m}(\mathbf{x}_i) \right]^2$$

- Classification – "0-1 loss"  $\Rightarrow$  misclassification error

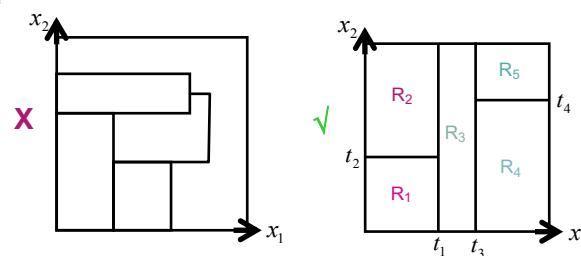
$$\{\hat{c}_m, \hat{R}_m\}_1^M = \arg \min_{\{c_m, R_m\}_1^M} \sum_{i=1}^N I\left( y_i \neq \sum_{m=1}^M c_m I_{R_m}(\mathbf{x}_i) \right)$$

## Tree Induction

### Theory Overview (2)

- Search Strategy
  - Unrestricted optimization with respect to  $\{R_m\}_1^M$  is difficult!
  - Region restrictions

- Disjoint
- Cover input space
- "Simple"



## Tree Induction

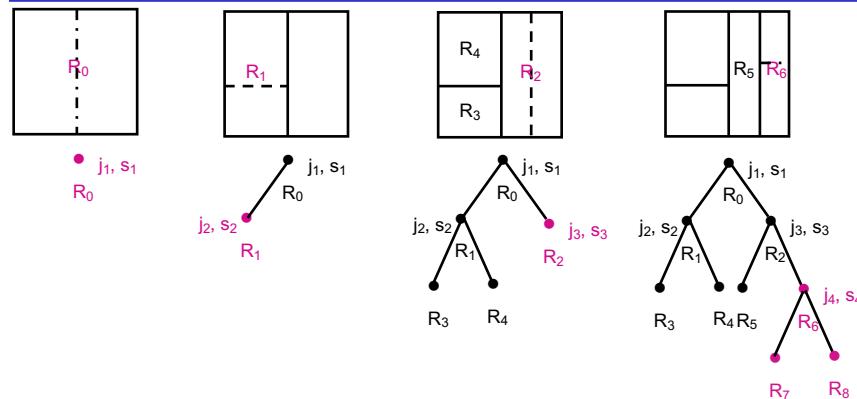
### Theory Overview (3)

- Search Strategy – "Simple" Regions
    - Let  $S_j$  be the set of all possible values of  $x_j$
    - Define "splits" –  $s_j \subseteq S_j$  – according to attribute type
      - Numeric (orderable):  $I(x_j \in s_j) = I(x_j \leq t_j)$
      - Categorical (nominal – unorderable): explicitly delineated
- $x_j = \text{occupation} \Rightarrow s_j = \{\text{manager, student}\}$
- Regions as "conjunctive" rules:  $x_1 \leq t_1 \text{ AND } x_2 \leq t_2 \Rightarrow R_1$

$$R = \bigcap_{j=1}^n s_j \Rightarrow I(\mathbf{x} \in R) = \prod_{j=1}^n I(x_j \in s_j)$$

## Tree Induction

### Binary Tree Connection



- Internal nodes represent splits
- Terminal nodes represent final regions defining model

## Tree Induction

### Growing Algorithm

- Greedy Iterative procedure

- Starting with a single region -- i.e., all given data
  - At the  $m$ -th iteration:

```
for each region R
  for each attribute  $x_j$  in R
    for each possible split  $s_j$  of  $x_j$ 
      record change in score when we partition  $R$  into  $R^l$  and  $R^r$ 
```

Choose  $(x_j, s_j)$  giving maximum improvement to fit

Replace  $R$  with  $R^l$ ; add  $R^r$

- When should we stop? i.e., how large should we grow the tree?

## Tree Induction

### Growing Algorithm (2)

- Split Scoring: *Regression*

- We seek split-attribute  $j$  and split-point  $s$  that solve

$$\min_{j,s} \left[ \min_{c_1} \sum_{\mathbf{x}_i \in R_{(j,s)}^l} (y_i - c_1)^2 + \min_{c_2} \sum_{\mathbf{x}_i \in R_{(j,s)}^r} (y_i - c_2)^2 \right]$$

where

$$R_{(j,s)}^l = \{\mathbf{x} \mid x_j \leq s\} \quad \text{and} \quad R_{(j,s)}^r = \{\mathbf{x} \mid x_j > s\}$$

- For any choice  $j$  and  $s$ , the inner minimization is solved by

$$\hat{c}_1 = \text{avg}\{y_i \mid \mathbf{x}_i \in R_{(j,s)}^l\} \quad \text{and} \quad \hat{c}_2 = \text{avg}\{y_i \mid \mathbf{x}_i \in R_{(j,s)}^r\}$$

## Tree Induction

### Growing Algorithm (3)

- Split Scoring: *Classification*

- In a node  $m$  representing a region  $R_m$  with  $N_m$  points, let

$$\hat{p}_{m,k} = \frac{1}{N_m} \sum_{\mathbf{x}_i \in R_m} I(y_i = k)$$

proportion of class  $k$  points in node  $m$

- We classify the points in node  $m$  to class  $\hat{c}_m = \arg \max_k \hat{p}_{m,k}$  – i.e., **majority class**

- $\hat{c}_m$  minimizes misclassification error:

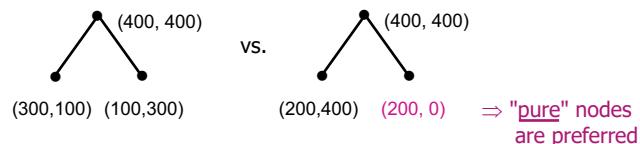
$$\begin{aligned} & \frac{1}{N_m} \sum_{\mathbf{x}_i \in R_m} I(y_i \neq \hat{c}_m) \\ &= 1 - \hat{p}_{m,\hat{c}_m} \end{aligned}$$

## Tree Induction

### Growing Algorithm (4)

- Split Scoring: Classification (Cont.)

- Problem:



Both splits produce an error of 0.25!

- Surrogate score criterion: "*Gini*" index of diversity

$$G(\hat{p}_{m,1}, \hat{p}_{m,2}, \dots, \hat{p}_{m,K}) = \sum_{k=1}^K \hat{p}_{m,k} (1 - \hat{p}_{m,k}) = 1 - \sum_{k=1}^K \hat{p}_{m,k}^2$$

- Same population solution!

- Max diversity (min purity):

$$G\left(\frac{1}{K}, \frac{1}{K}, \dots, \frac{1}{K}\right) = 1 - \frac{1}{K}$$

- Min diversity (max purity):

$$G(0,0,\dots,1,0,\dots,0) = 0$$

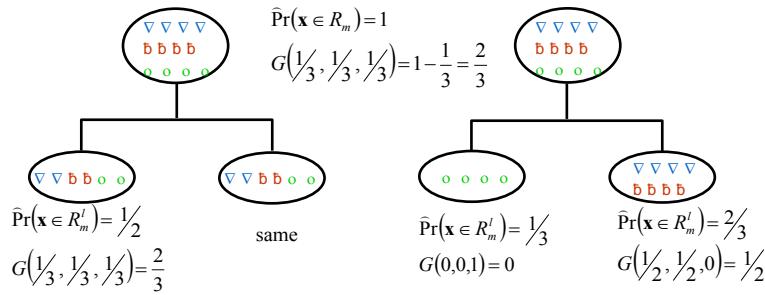
## Tree Induction

### Growing Algorithm (5)

- Split Scoring: Classification (Cont.)

- Goodness of split: difference between purity of parent and weighted sum of purity of daughters

$$\hat{I}_m(j, s_{jm}) = \hat{\Pr}(\mathbf{x} \in R_m) \cdot G(R_m) - \hat{\Pr}(\mathbf{x} \in R'_m) \cdot G(R'_m) - \hat{\Pr}(\mathbf{x} \in R''_m) \cdot G(R''_m)$$



## Tree Induction

### Growing Algorithm (6)

- Split Scoring: Classification (Cont.)

- Gini index not the only "diversity" measure

- Second order entropy:  $H^{(2)} = -\sum_{k=1}^K \hat{p}_k^2 (\mathcal{Y} = c_k \mid \mathbf{x} \in R)$

- Has some bias towards “equal” size splits

- Ordinary entropy:  $H = -\sum_{k=1}^K \hat{p}_k \log \hat{p}_k$

- Used in C4.5

- Gets unhappy when  $\hat{p}_k \approx 0$

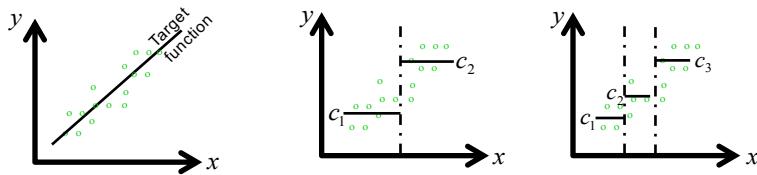
- Any super-linear function of  $\{\hat{p}_k\}_{k=1}^K$  will work  $\Rightarrow$  Try them out!

## Tree Induction

### Pruning

- Dilemma

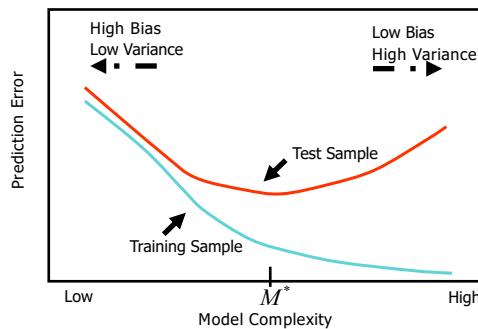
- If tree (# of regions) is too small, then piecewise constant approximation is too crude (**bias**)  $\Rightarrow$  increased errors
- If tree is too large, then it fits the training data too closely (overfitting, increased **variance**)  $\Rightarrow$  increased errors



## Tree Induction

### Pruning (2)

- Bias-Variance tradeoff



- Right sized tree,  $M^*$ , when test error is at a minimum
- Error on the training is not a useful estimator!
  - If test set is not available, need alternative method (e.g., Cross-validation)

## Tree Induction

### Pruning (3)

- Two strategies

– *Prepruning* - stop growing a branch when information becomes unreliable

- $\#(R_m)$  – i.e., number of data points, too small

⇒ same bound everywhere in the tree

- Next split not worthwhile

⇒ Not sufficient condition

- *Postpruning* - take a fully-grown tree and discard unreliable parts (i.e., not supported by test data)

- C4.5: pessimistic pruning

- CART: cost-complexity pruning (*more statistically grounded*)

## Key Features

### ■ Ability to deal with irrelevant inputs

- i.e., automatic variable subset selection

- Measure anything you can measure

- Score provided for selected variables ("importance")

### ■ No data preprocessing needed

- Naturally handle all types of variables

- numeric, binary, categorical

- Invariant under monotone transformations:  $x_j = g_j(x_j)$

- Variable scales are irrelevant

- Immune to bad  $x_j$  - distributions (e.g., outliers)

## Key Features (2)

- Computational scalability
  - Relatively fast:  $O(nN \log N)$
- Missing value tolerant
  - Moderate loss of accuracy due to missing values
  - Handling via "surrogate" splits
- "Off-the-shelf" procedure
  - Few tunable parameters
- Interpretable model representation
  - Binary tree graphic

## Surrogate Splits and Their Uses

### Missing Predictor Values

- Suppose  $\mathbf{x} = \{x_1, x_2, \dots, \underset{\text{missing value}}{x_l}, \dots, x_n\}$
- Assume we have a split  $I(x_l \in S)$  in our tree
  - What to do if  $x_l = NA$ ?
    - Could delete cases  $\Rightarrow$  data depletion
    - Don't go any deeper (do with what tree has learned so far for this case)
    - Allow case to follow majority (assumes all missing are somehow typical)
    - Allow missing to be a separate value (missing are indistinguishable)
    - Try to fill in (*impute*) value

## Surrogate Splits and Their Uses

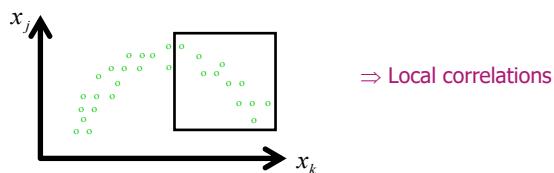
### Missing Predictor Values (2)

- A more general approach
  - When considering predictor  $x_j$  for a split, only  $\{x_i \mid x_{ij} \neq \text{missing}\}$  are used
    - A penalty is added for number of missing
  - Suppose optimal, *primary*, split of  $R_m = (x_{j(m)}, s_m)$ 
    - Idea: use other var's to predict this split
  - Consider “surrogate” response  $\tilde{y} = I(x_{j(m)} \in s_m)$
  - For each var  $x_k$  (different from  $x_{j(m)}$ )
    - Compute  $c_k = \max_s \text{corr}[\tilde{y}, I(x_k \in s)]$  (“agreement” score)
    - $s_k$  maximizing value
    - Order  $\{k \neq j(m)\}$  on decreasing  $c_k \Rightarrow 1^{\text{st}}$  surrogate,  $2^{\text{nd}}$  surrogate, ...

## Surrogate Splits and Their Uses

### Missing Predictor Values (3)

- Surrogates exploit linear and nonlinear relationships among the input variables



i.e., Redundant var's can be helpful!

- Competitors vs. Surrogates

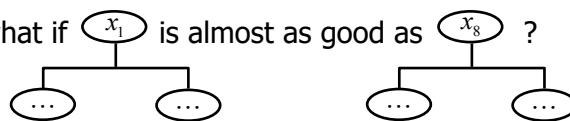
Class A	100
Class B	100
Class C	100

	Primary		Competitor		Surrogate	
	L	R	L	R	L	R
Class A	90	10	80	20	78	22
Class B	80	20	25	75	74	26
Class C	15	85	14	86	21	79

## Surrogate Splits and Their Uses

### Ranking Predictor Variables

- Which input variables are the most important?
  - i.e., relative influence or contribution in predicting the response
  - Higher importance variables are more likely to be of interest
- Variable masking
  - How to rank those vars that, while not giving the best split, may give the second, third, etc.?
  - i.e., what if  $x_1$  is almost as good as  $x_8$  ?



## Surrogate Splits and Their Uses

### Ranking Predictor Variables (2)

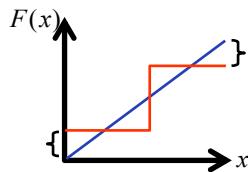
- Measure of relevance

$$\mathfrak{I}(x_l) = \sum_{t \in T} \hat{\mathfrak{I}}(v(t), \tilde{s}_{v(t)}) \cdot I(v(t) = l)$$

- i.e., sum the "goodness of split" scores whenever  $x_l$  is used in a surrogate split
- Sum is over all internal nodes in the tree
  - If  $x_l$  was used as the primary split in some node, then corresponding  $\hat{\mathfrak{I}}(l, s_l)$  is used in the summation
  - We should have  $\mathfrak{I}(x_1) \approx \mathfrak{I}(x_8)$  in previous scenario

## Tree Limitations

- Discontinuous piecewise constant model



- In order to have many splits you need to have a lot of data
  - In high-dimensions, you often run out of data after a few splits
- Also note error is bigger near region boundaries

## Tree Limitations (2)

- Not good for low interaction  $F^*(\mathbf{x})$

– e.g.,  $F^*(\mathbf{x}) = a_o + \sum_{j=1}^n a_j x_j$  is worst function for trees

$$= \sum_{j=1}^n f_j^*(x_j) \quad (\text{no interaction, additive})$$

- In order for  $x_j$  to enter model, must split on it
  - Path from root to node is a product of indicators

- Not good for  $F^*(\mathbf{x})$  that has dependence on many variables

- Each split reduces training data for subsequent splits (data fragmentation)

## Tree Limitations (3)

- High variance caused by greedy search strategy (local optima)
  - Errors in upper splits are propagated down to affect all splits below it
    - ⇒ Small changes in data (sampling fluctuations) can cause big changes in tree
    - Very deep trees might be questionable
    - Pruning is important

## What to do next?

- Live with problems
- Use other methods (when possible)
- Fix-up trees: use "ensembles"
  - Maintain advantages while dramatically increasing accuracy